

```
In [19]: # check current directory
pwd()
"/Users/xyliu/Desktop/ISE 616 final/ISE-616-final"
```

```
In [40]: using JuMP
using Ipopt
using LinearAlgebra
using MathOptInterface
const MOI = MathOptInterface
using Random
```

# Modeling

## function of one-hot encoding

```
In [21]: #####
# Encoding information type
#####

"""
    ZGEncodingInfo

Stores the information needed to convert between
original (integer-coded) categorical/group features
and their one-hot encodings.

Fields:
- k_z::Vector{Int}      : number of categories for each categorical component
- z_start::Vector{Int}   : start index (column) of each component in Z_oneh
- num_g::Int             : total number of groups
"""

struct ZGEncodingInfo
    k_z::Vector{Int}
    z_start::Vector{Int}
    num_g::Int
end

#####
# Original (Z, group) → Reduced (Z_enc, G_enc)
#####

"""
    encode_zg_reduced(Z::AbstractMatrix{<:Integer},
                      group::AbstractVector{<:Integer})

Encode categorical features Z and group indexes into a reduced dummy
representation:

- Each categorical component l with k_l levels is mapped to (k_l - 1)
  dummy variables. The last category k_l is the baseline (all zeros).

```

- The group variable with `num_g` levels is mapped to  $(\text{num\_g} - 1)$  dummy variables. The last group `num_g` is the baseline (all zeros).

**Input:**

- `Z` :  $N \times m$  matrix, each column `l` stores an integer in  $\{1, \dots, k_l\}$
- `group` : length- $N$  vector, each entry in  $\{1, \dots, \text{num\_g}\}$

**Output:**

- `Z_enc` :  $N \times \sum_l (k_l - 1)$  matrix (reduced dummies for `Z`)
- `G_enc` :  $N \times (\text{num\_g} - 1)$  matrix (reduced dummies for `group`)
- `info` : `ZGEncodingInfo`, stores the encoding scheme for later decoding

```

function encode_zg_reduced(Z::AbstractMatrix{<:Integer},
                           group::AbstractVector{<:Integer})
    N, m = size(Z)
    @assert length(group) == N "group must have length N"

    #  $k_{z[l]}$  = number of categories for component  $l$  (assumed  $\{1, \dots, k_l\}$ )
    k_z = [maximum(Z[:, l]) for l in 1:m]

    # Column start indices in the reduced dummy matrix Z_enc.
    # Component  $l$  uses  $(k_{z[l]} - 1)$  columns (baseline = category  $k_{z[l]}$ ).
    z_start = Vector{Int}(undef, m)
    col = 1
    for l in 1:m
        z_start[l] = col
        d_l = max(k_z[l] - 1, 0) # number of dummies for component  $l$ 
        col += d_l
    end
    total_z = col - 1 # total number of columns in Z_enc

    # Allocate Z_enc:  $N \times \text{total}_z$ 
    Z_enc = zeros(Int8, N, total_z)

    for i in 1:N
        for l in 1:m
            val = Z[i, l]
            @assert 1 ≤ val ≤ k_z[l] "Category out of range in column $l"

            k_l = k_z[l]
            d_l = max(k_l - 1, 0)

            # If  $d_l == 0$ , there is effectively a single category (no col
            if d_l > 0 && val < k_l
                col0 = z_start[l] # start index for component  $l$ 
                Z_enc[i, col0 + val - 1] = 1 # categories  $1..(k_l-1)$ 
            end
            # If  $val == k_l$ , baseline category -> all zeros for this component
        end
    end

    # Encode group using  $(\text{num\_g} - 1)$  dummies; last group is baseline.
    num_g = maximum(group)
    d_g = max(num_g - 1, 0)
    G_enc = zeros(Int8, N, d_g)

    if d_g > 0
        for i in 1:N
            gi = group[i]
            @assert 1 ≤ gi ≤ num_g "Group index out of range at row $i"
    
```

```

        if gi < num_g
            # Non-baseline group in {1, ..., num_g - 1}
            G_enc[i, gi] = 1
        else
            # Baseline group num_g -> all zeros
        end
    end
end

info = ZGEncodingInfo(k_z, z_start, num_g)
return Z_enc, G_enc, info
end

```

encode\_zg\_reduced

## test one-hot encode

```

In [22]: # testing
Z = [
    1 1;    # row 1
    2 3;    # row 2
    3 2;    # row 3
    1 3;    # row 4
]
# group: length-N vector, values in {1,2,3}
group = [1, 2, 1, 3]

#####
# 2. Run encoding
#####

Z_enc, G_onehot, info = encode_zg_reduced(Z, group)
println("Original Z:")
println(Z)
println()

println("Encoded Z_enc (reduced one-hot):")
println(Z_enc)
println("size(Z_enc) = ", size(Z_enc))
println()

println("Group one-hot G_onehot:")
println(G_onehot)
println("size(G_onehot) = ", size(G_onehot))
println()

println("Encoding info:")
println("  k_z      = ", info.k_z)
println("  z_start  = ", info.z_start)
println("  num_g   = ", info.num_g)

```

Original Z:  
[1 1; 2 3; 3 2; 1 3]

Encoded Z\_enc (reduced one-hot):  
Int8[1 0 1 0; 0 1 0 0; 0 0 0 1; 1 0 0 0]  
size(Z\_enc) = (4, 4)

Group one-hot G\_onehot:  
Int8[1 0; 0 1; 1 0; 0 0]  
size(G\_onehot) = (4, 2)

Encoding info:  
k\_z = [3, 3]  
z\_start = [1, 3]  
num\_g = 3

In [23]: #####  
# Reduced (Z\_enc, G\_enc) → Original (Z, group)  
#####

.....  
decode\_zg\_reduced(Z\_enc::AbstractMatrix{<:Integer},  
G\_enc::AbstractMatrix{<:Integer},  
info::ZGEncodingInfo)

Decode reduced dummy encoded categorical and group features back to their original integer-coded form.

For each categorical component l with k\_l levels:

- If the corresponding block has all zeros, we recover category k\_l (the baseline).
- Otherwise, the position of the 1 determines the category in {1, ..., k\_l}

For the group variable with num\_g levels:

- If the row in G\_enc is all zeros, we recover group num\_g (baseline).
- Otherwise, the position of the 1 determines the group in {1, ..., num\_g}

**Input:**

- Z\_enc : N × sum\_l (k\_l - 1) matrix (reduced dummies for Z)
- G\_enc : N × (num\_g - 1) matrix (reduced dummies for group)
- info : ZGEncodingInfo produced earlier by `encode\_zg\_reduced`

**Output:**

- Z : N × m matrix, each entry in {1, ..., k\_l}
- group : length-N vector, each entry in {1, ..., num\_g}

.....

```
function decode_zg_reduced(Z_enc::AbstractMatrix{<:Integer},
                           G_enc::AbstractMatrix{<:Integer},
                           info::ZGEncodingInfo)
    N, total_z = size(Z_enc)
    m = length(info.k_z)

    # Check group encoding consistency
    num_g = info.num_g
    d_g = max(num_g - 1, 0)
    @assert size(G_enc, 1) == N "Z_enc and G_enc must have same number of
    @assert size(G_enc, 2) == d_g "G_enc columns must be num_g - 1"

    # Recover categorical matrix Z (N × m)
    Z = zeros(Int, N, m)
```

```

for l in 1:m
    k_l = info.k_z[l]
    d_l = max(k_l - 1, 0)
    s = info.z_start[l]
    e = s + d_l - 1

    if d_l == 0
        # Only one category exists for this component; always categor
        for i in 1:N
            Z[i, l] = 1
        end
        continue
    end

    @assert e ≤ total_z "Z_enc has too few columns for component $l"

    # View the reduced dummy block for component l
    sub = @view Z_enc[:, s:e] # N × d_l
    for i in 1:N
        idx = findfirst==(1), sub[i, :])
        if idx === nothing
            # All zeros: baseline category k_l
            Z[i, l] = k_l
        else
            # Non-baseline category in {1, ..., k_l - 1}
            Z[i, l] = idx
        end
    end
end

# Recover group vector (length N) from reduced dummies
group = Vector{Int}(undef, N)
if d_g == 0
    # Only one group; always group 1
    for i in 1:N
        group[i] = 1
    end
else
    for i in 1:N
        idx = findfirst==(1), G_enc[i, :])
        if idx === nothing
            # Baseline group num_g
            group[i] = num_g
        else
            # Non-baseline group in {1, ..., num_g - 1}
            group[i] = idx
        end
    end
end

return Z, group
end

```

decode\_zg\_reduced

## test one-hot decode

In [24]:

```
Z, group = decode_zg_reduced(Z_enc, G_onehot, info)
Z, group
```

```
([1 1; 2 3; 3 2; 1 3], [1, 2, 1, 3])
```

## DAG builder

In [25]:

```
#####
# Graph structure for sample-level DAGs
# compatible with Theorem 2 (graph-based formulation)
#####

# -----
# Arc kind (structure only)
# -----



#####
ArcKind

Abstract type for different kinds of arcs in the DAG.
We distinguish between:
- `CatArc` : categorical feature transitions
- `TermArc` : terminal transitions to (m+1, 0) with a chosen group g
#####
abstract type ArcKind end

#####


CatArc

Categorical arc corresponding to choosing category `c` for component `k`.

Fields:
- k      : which categorical component (1..m)
- c      : chosen category in {1, ..., k_z[k]}
- d_prev : previous accumulated categorical distance at state (k-1, d_prev)
- d      : new accumulated distance at state (k, d)
#####


struct CatArc <: ArcKind
    k::Int
    c::Int
    d_prev::Float64
    d::Float64
end

#####


TermArc

Terminal arc from (m, d) to (m+1, 0) with a chosen destination group g.

Fields:
- d : accumulated categorical distance at state (m, d)
- g : destination group in {1, ..., num_g}
#####
struct TermArc <: ArcKind
    d::Float64
    g::Int
end

# -----
# Arc and per-sample DAG
# -----
```

```

#####
Arc

Directed edge in the DAG for a single sample i.

Fields:
- src : index of the source node in `nodes`
- dst : index of the destination node in `nodes`
- kind : arc type (CatArc or TermArc), which encodes all information
          needed to build  $w^i(e; \beta, \lambda, r_i, \dots)$  later
#####
struct Arc
    src::Int
    dst::Int
    kind::ArcKind
end

#####

SampleDAG

Graph structure associated with a single sample i, compatible with
Theorem 2 (graph-based formulation).

Fields:
- sample_index : index i of the sample this DAG corresponds to
- nodes : vector of DP states (k, d), plus the terminal node (m+1,
           each node is a Tuple{Int,Float64})
- arcs : list of directed edges with their arc kind (no numeric w
- source : index of the source node (corresponds to state (0, 0.0))
- sink : index of the terminal node (corresponds to state (m+1, 0
#####
struct SampleDAG
    sample_index::Int
    nodes::Vector{Tuple{Int,Float64}}
    arcs::Vector{Arc}
    source::Int
    sink::Int
end

# -----
# Categorical encoding info
# -----
#####

CatEncodingInfo

Encoding information for categorical features (structure level only).

Fields:
- k_z::Vector{Int} : number of categories per component (length m).
                      For component k, categories are {1, ..., k_z[k]}.
- num_g::Int       : total number of groups (used for terminal arcs).

```

```

.....
struct CatEncodingInfo
    k_z::Vector{Int}
    num_g::Int
end

# -----
# Build DAG structure for one sample i
# -----
.....
    build_sample_dag_structure(i, info, delta, z_i) -> SampleDAG

Build the DAG structure  $G^i = (V^i, A^i)$  for a fixed sample  $i$ ,  

compatible with Theorem 2 (graph-based formulation), using our setup.

This function ONLY builds:
- the nodes (states  $(k, d)$  plus the terminal  $(m+1, 0)$ ),
- the arcs with their structural information (CatArc or TermArc).

It does NOT compute numeric weights  $w^i(e)$ . Those should be constructed  

later as expressions of  $(\beta, \lambda, r_i, y_i, g_i, B_{\{g_i\}}, C_{\{g_i\}}, \dots)$ .

Arguments:
- i : sample index
- info : CatEncodingInfo (k_z, num_g)
- delta : length-m vector of  $\delta_k$  in
           $d(z, z^i) = \sum_k \delta_k * 1[z_k \neq z_{k^i}]$ 
- z_i : length-m vector of original categories for sample i,
        each  $z_i[k] \in \{1, \dots, k_z[k]\}$ 

Output:
- SampleDAG describing the structure of  $G^i$ 
.....
function build_sample_dag_structure(
    i::Int,
    info::CatEncodingInfo,
    delta::AbstractVector{<:Real},
    z_i::AbstractVector{<:Integer}
) :: SampleDAG
    k_z = info.k_z
    num_g = info.num_g
    m = length(k_z)

    # -----
    # 0. Sanity checks
    # -----
    @assert length(delta) == m "delta must have length m"
    @assert length(z_i) == m "z_i must have length m"

    # -----
    # 1. Enumerate states (k, d) with per-layer dedup
    # -----
    nodes = Vector{Tuple{Int,Float64}}()
    node_index = Dict{Tuple{Int,Float64},Int}()

    # Source state (0, 0.0)
    push!(nodes, (0, 0.0))
    node_index[(0, 0.0)] = 1
    source_idx = 1

```

```

arcs = Vector{Arc}()

# current_layer holds all unique states (k-1, d_prev)
current_layer = [(0, 0.0)]

for k in 1:m
    next_layer = Tuple{Int,Float64}[]
    seen_next = Set{Tuple{Int,Float64}}()

    k_l = k_z[k]
    δ_k = float(delta[k])
    z_i_k = z_i[k]

    for (k_prev, d_prev) in current_layer
        @assert k_prev == k - 1

        # Enumerate all categories c ∈ {1, ..., k_l}
        for c in 1:k_l
            mismatch = (c != z_i_k)
            d = d_prev + (mismatch ? δ_k : 0.0)
            state = (k, d)

            # Add state to global node list if new
            if !haskey(node_index, state)
                push!(nodes, state)
                node_index[state] = length(nodes)
            end

            # Ensure each (k, d) appears at most once in next_layer
            if !(state in seen_next)
                push!(next_layer, state)
                push!(seen_next, state)
            end

            # Add categorical arc from (k-1, d_prev) to (k, d)
            src = node_index[(k-1, d_prev)]
            dst = node_index[state]
            kind = CatArc(k, c, d_prev, d)
            push!(arcs, Arc(src, dst, kind))
        end
    end

    # Move to next layer (already deduplicated)
    current_layer = next_layer
end

# States with k = m are the final DP layer S_1^i
# current_layer is already deduplicated, but we keep the name for cla
S1_states = current_layer

# -----
# 2. Add terminal node (m+1, 0.0) and terminal arcs
# -----
terminal_state = (m+1, 0.0)
push!(nodes, terminal_state)
node_index[terminal_state] = length(nodes)
sink_idx = node_index[terminal_state]

# For each (m, d) in S_1^i, and for each group g, add a TermArc

```

```

    for (k_state, d) in S1_states
        @assert k_state == m
        src = node_index[(m, d)]
        for g in 1:num_g
            kind = TermArc(d, g)
            push!(arcs, Arc(src, sink_idx, kind))
        end
    end

    # -----
    # 3. Return SampleDAG
    # -----
    return SampleDAG(i, nodes, arcs, source_idx, sink_idx)
end

```

build\_sample\_dag\_structure

## test DAG builder

In [26]:

```

i = 2001

k_z = [2, 2]      # each component 2 cats
delta = [1.0, 1.0] # δ₁ = δ₂ = 1
z_i = [1, 2]      # sample i: z₁ = 1, z₂ = 2
num_g = 2

info = CatEncodingInfo(k_z, num_g)

dag_i = build_sample_dag_structure(i, info, delta, z_i)
dag_i

```

```

SampleDAG(2001, [(0, 0.0), (1, 0.0), (1, 1.0), (2, 1.0), (2, 0.0), (2, 2.0),
(3, 0.0)], Arc[Arc(1, 2, CatArc(1, 1, 0.0, 0.0)), Arc(1, 3, CatArc(1, 2, 0.0, 1.0)),
Arc(2, 4, CatArc(2, 1, 0.0, 1.0)), Arc(2, 5, CatArc(2, 2, 0.0, 0.0)), Arc(3, 6, CatArc(2, 1, 1.0, 2.0)),
Arc(3, 4, CatArc(2, 2, 1.0, 1.0)), Arc(4, 7, TermArc(1.0, 1)), Arc(4, 7, TermArc(1.0, 2)), Arc(5, 7, TermArc(0.0, 1)),
Arc(5, 7, TermArc(0.0, 2)), Arc(6, 7, TermArc(2.0, 1)), Arc(6, 7, TermArc(2.0, 2))], 1, 7)

```

## modeling

In [27]:

```

using JuMP

"""
    build_group_dro_graph_model(
        X, Z, group, y,
        encinfo,
        delta,
        A_group,
        B_group, C_group,
        gamma_x,
        ε,
        optimizer
    ) -> (model, meta)

```

Build our group-dependent, graph-based DRO logistic regression model, using the same reduced encoding convention as `ZGEncodingInfo` / `encode\_zg\_reduced`.

**Arguments**

-----

- X ::  $N \times n_x$  matrix of continuous features.
- Z ::  $N \times m$  matrix of original categorical features.  
Entry  $Z[i, k] \in \{1, \dots, k_z[k]\}$  (no one-hot).
- group :: length- $N$  vector of group indices  $g_i \in \{1, \dots, num_g\}$ .
- y :: length- $N$  vector of labels in  $\{-1, +1\}$ .
- encinfo :: ZGEncodingInfo  
( $k_z, z_{start}, num_g$ ) describing reduced encoding blocks  
for the categorical features and groups.
- delta :: length- $m$  vector  $\delta_k$  used in  
 $d_{cat}(z, z^i) = \sum_k \delta_k * 1[z_k \neq z_k^i]$ .
- A\_group :: length- $num_g$  vector  $A_g$   
continuous-part metric weights in  
 $A_g \sum_j \gamma_j |x_j - x_j^i|$ .
- B\_group :: length- $num_g$  vector  $B_g$   
categorical-part metric weight.
- C\_group :: length- $num_g$  vector  $C_g$   
group-change penalty.
- gamma\_x :: length  $n_x$  vector  $\gamma_j$  for continuous features.  
Continuous dual constraint will be  
 $|\beta_{xj}| \leq \lambda * A_{min} * \gamma_j$ ,  
where  $A_{min} = \text{minimum}(A_{group})$ .
- $\varepsilon$  :: Wasserstein radius  $\varepsilon$ .
- optimizer: optimizer constructor for JuMP, e.g.  
`optimizer_with_attributes(Mosek.Optimizer, "QUIET" => true)`

**Returns**

-----

- model :: JuMP.Model
- meta :: NamedTuple: (encinfo=encinfo, dags=dags, n\_nodes=n\_nodes)

```
"""
function build_group_dro_graph_model(
    X::AbstractMatrix{<:Real},
    Z::AbstractMatrix{<:Integer},
    group::AbstractVector{<:Integer},
    y::AbstractVector{<:Integer},
    encinfo::ZGEncodingInfo,
    delta::AbstractVector{<:Real},
    A_group::AbstractVector{<:Real},
    B_group::AbstractVector{<:Real},
    C_group::AbstractVector{<:Real},
    gamma_x::AbstractVector{<:Real},
    ε::Real,
    optimizer,
)
    # -----
    # 0. Dimensions & sanity checks
    # -----
    N, n_x = size(X)
```

```

N_Z, m = size(Z)
@assert N_Z == N "X and Z must have the same number of rows (samples)"
@assert length(group) == N "group must have length N."
@assert length(y) == N "y must have length N."
@assert length(delta) == m "delta must have length m."
@assert length(gamma_x) == n_x "gamma_x must have length n_x."

k_z      = encinfo.k_z
z_start = encinfo.z_start
num_g   = encinfo.num_g

@assert length(k_z) == m "encinfo.k_z must have length m."
@assert length(z_start) == m "encinfo.z_start must have length m."
@assert length(A_group) == num_g "A_group length must equal num_g."
@assert length(B_group) == num_g "B_group length must equal num_g."
@assert length(C_group) == num_g "C_group length must equal num_g."

# Total length of  $\beta_z$  under reduced encoding:
# last block starts at z_start[m], length (k_z[m] - 1)
# so p_z = z_start[m] + (k_z[m] - 1) - 1
p_z = z_start[end] + (k_z[end] - 1) - 1

# Small positive constant for log(exp(inner) - 1) domain
# We will enforce: r[i] +  $\lambda$  * inner_coeff  $\geq \eta$ 
η = 1e-6

# -----
# 1. Build per-sample DAGs (discrete part only)
# -----
info = CatEncodingInfo(k_z, num_g)

dags      = Vector{SampleDAG}(undef, N)
n_nodes   = Vector{Int}(undef, N)
max_nodes = 0

for i in 1:N
    # Z[i, :] is an AbstractVector{<:Integer}
    dags[i] = build_sample_dag_structure(i, info, delta, Z[i, :])
    n_nodes[i] = length(dags[i].nodes)
    max_nodes = max(max_nodes, n_nodes[i])
end

# -----
# 2. Create JuMP model & decision variables
# -----
model = Model(optimizer)

# Wasserstein dual variable  $\lambda \geq 0$ 
@variable(model, λ >= 0.0)

# Per-sample slack variables  $r_i \in \mathbb{R}$  (free),
# domain of  $\log(\exp(\dots)-1)$  will be enforced via extra constraints
@variable(model, r[1:N])

# Logistic regression parameters
@variable(model, β₀)                      # intercept
@variable(model, β_x[1:n_x])                # continuous coefficients
@variable(model, β_z[1:p_z])                # categorical coefficients (reduced)
@variable(model, β_grp[1:(num_g-1)])# group coefficients (reduced)

```

```

# Graph dual variables  $\mu_{i,v}$  for each sample  $i$  and each node  $v$ 
@variable(model, mu[1:N, 1:max_nodes])

# -----
# 3. Objective:  $\lambda \varepsilon + (1/N) \sum r_i$ 
# -----
@objective(model, Min, lambda * epsilon + (1.0 / N) * sum(r[i] for i in 1:N))

# -----
# 4. Continuous part dual constraints with A_group
#
# Metric for  $x$  uses A_g:
#  $d_x(x^i, x; g) = A_g \sum_j \gamma_j |x_j - x_j^i|$ 
# Dual boundedness  $\Rightarrow$  for each  $j$ :
#  $|\beta_{\{xj\}}| \leq \lambda * \gamma_j * \min_g A_g$ 
# We enforce:
#  $-\lambda * A_{\min} * \gamma_j \leq \beta_{x[j]} \leq \lambda * A_{\min} * \gamma_j$ 
# -----
A_min = minimum(A_group)

for j in 1:n_x
    @constraint(model, beta_x[j] <= lambda * A_min * gamma_x[j])
    @constraint(model, -beta_x[j] <= lambda * A_min * gamma_x[j])
end

# -----
# 5. Outer logistic inequality:
#  $y^i (\beta_0 + \beta_x^T x^i) \geq -\mu_i(0,0) + \mu_i(m+1,0)$ 
# -----
for i in 1:N
    dag = dags[i]
    s = dag.source
    t = dag.sink

    # Left-hand side:  $y_i * (\beta_0 + \beta_x^T x^i)$ 
    lhs = y[i] * (beta_0 + sum(beta_x[j] * X[i, j] for j in 1:n_x))

    # Right-hand side:  $-\mu_i(\text{source}) + \mu_i(\text{sink})$ 
    rhs = -mu[i, s] + mu[i, t]

    @constraint(model, lhs >= rhs)
end

# -----
# 6. Edge constraints:  $\mu_t(e) - \mu_s(e) \geq w^i(e)$ 
#
# - For CatArc( $k, c, \dots$ ):
#    $w^i(e) = -y^i \beta_{\{z_k\}^T z_k(c)}$ 
#   Reduced encoding aligned with ZGEncodingInfo:
#   if  $c < k_z[k]$ ,  $\beta_{\{z_k\}^T z_k(c)} = \beta_z[z_start[k] + c - 1]$ 
#   if  $c = k_z[k]$ , baseline  $\Rightarrow 0$ .
#
# - For TermArc( $d, g$ ):
#    $w^i(e) =$ 
#    $-y^i \beta_{grp}^T \phi_g(g)$ 
#    $-log(exp(r_i + \lambda (B_{\{g_i\}} d + C_{\{g_i\}} 1[g \neq g_i])) - 1)$ 
#
# with group reduced encoding:
#   if  $g < num_g$ :  $\beta_{grp}^T \phi_g(g) = \beta_{grp}[g]$ 
#   if  $g = num_g$ : baseline  $\Rightarrow 0$ .

```

```

#
#   Additionally, domain constraints:
#        $r_i + \lambda (B_{\{g_i\}} d + C_{\{g_i\}} 1[g \neq g_i]) \geq \eta$ 
#   ensure that  $\log(\exp(\text{inner}) - 1)$  is well-defined.
# ----

for i in 1:N
    dag = dags[i]
    y_i = y[i]
    g_i = group[i]

    for arc in dag.arcs
        src = arc.src
        dst = arc.dst

        if arc.kind isa CatArc
            kind = arc.kind::CatArc
            k = kind.k
            c = kind.c
            k_l = k_z[k]

            # Categorical part:
            #  $w_{cat}^i(e) = -y_i \beta_{\{z_k\}}^T z_k(c)$ 
            if c < k_l
                idx = z_start[k] + (c - 1)
                w_expr = -y_i * beta_z[idx]
            else
                w_expr = 0.0
            end

            @constraint(model, mu[i, dst] - mu[i, src] >= w_expr)

        elseif arc.kind isa TermArc
            kind = arc.kind::TermArc
            d_val = kind.d
            g_choice = kind.g

            # Metric coefficient for categorical + group part:
            B_gi = float(B_group[g_i])
            C_gi = float(C_group[g_i])
            cross = (g_choice != g_i) ? C_gi : 0.0
            inner_coeff = B_gi * d_val + cross

            # --- Domain constraint:  $inner = r[i] + \lambda * inner_coeff \geq \eta$ 
            @NLconstraint(model, r[i] + lambda * inner_coeff >= eta)

            # ---  $w^i(e)$  constraint with nonlinear log/exp ---
            if g_choice < num_g
                # group linear term:  $-y_i * \beta_{grp}[g\_choice]$ 
                @NLconstraint(model,
                    mu[i, dst] - mu[i, src] >=
                    -y_i * beta_grp[g_choice] -
                    log(exp(r[i] + lambda * inner_coeff) - 1)
                )
            else
                # baseline group: no  $\beta_{grp}$  contribution
                @NLconstraint(model,
                    mu[i, dst] - mu[i, src] >=
                    -log(exp(r[i] + lambda * inner_coeff) - 1)
                )
            end
        end
    end
end

```

```

        else
            error("Unknown arc kind in DAG.")
        end
    end
end

# -----
# 7. (Optional) Give Ipopt a safe starting point inside the domain
# -----
set_start_value(λ, 1.0)
for i in 1:N
    set_start_value(r[i], 1.0)
end

# -----
# 8. Return model + metadata
# -----
meta = (
    encinfo = encinfo,
    dags = dags,
    n_nodes = n_nodes,

    β₀ = β₀,
    β_x = β_x,
    β_z = β_z,
    β_grp = β_grp,
    λ = λ,
    r = r,
    μ = μ,
)
return model, meta
end

```

build\_group\_dro\_graph\_model

## Toy example

```
In [28]: # ===== Toy data =====
N = 2      # two samples
n_x = 1     # one continuous feature
m = 2       # two categorical components
num_g = 2   # two groups

# Continuous features X: N × n_x
X = [
    0.0;
    1.0
]
X = reshape(X, N, n_x)  # 2×1 matrix

# Categorical features Z: N × m, z_{i,k} ∈ {1,2}
Z = [
    1 1;  # sample 1: (1,1)
    2 2   # sample 2: (2,2)
]

# Group indices g_i ∈ {1,2}
group = [1, 2]
```

```
# Labels  $y_i \in \{-1, +1\}$ 
y = [1, -1]

# Hamming weights  $\delta_k$ 
delta = [1.0, 1.0]

# Metric weights
A_group = [1.0, 2.0]    # continuous part weights A_g
B_group = [1.0, 1.0]    # categorical part weights B_g
C_group = [0.5, 0.5]    # group-change penalty C_g

# Continuous scaling  $\gamma_x$ 
gamma_x = [1.0]          # length n_x

# Wasserstein radius
ε = 0.1
```

0.1

In [29]:

```
# ===== Encoding info =====
k_z = [2, 2]

# reduced encoding block starts for categorical β_z
# block 1 starts at 1, length (k_z[1]-1) = 1
# block 2 starts at 2, length (k_z[2]-1) = 1
z_start = [1, 2]

struct ZGEncodingInfo
    k_z::Vector{Int}
    z_start::Vector{Int}
    num_g::Int
end

encinfo = ZGEncodingInfo(k_z, z_start, num_g)
```

ZGEncodingInfo([2, 2], [1, 2], 2)

In [30]:

```
# ===== Build model =====
model, meta = build_group_dro_graph_model(
    X,
    Z,
    group,
    y,
    encinfo,
    delta,
    A_group,
    B_group,
    C_group,
    gamma_x,
    ε,
    optimizer_with_attributes(Ipopt.Optimizer) # or just Ipopt.Optimizer
)

println("Model successfully built.")
println(model)
```

```

Model successfully built.
Min 0.1 λ + 0.5 r[1] + 0.5 r[2]
Subject to
β₀ + μ[1,1] - μ[1,7] ≥ 0
-β₀ - β_x[1] + μ[2,1] - μ[2,7] ≥ 0
β_z[1] - μ[1,1] + μ[1,2] ≥ 0
-μ[1,1] + μ[1,3] ≥ 0
β_z[2] - μ[1,2] + μ[1,4] ≥ 0
-μ[1,2] + μ[1,5] ≥ 0
β_z[2] - μ[1,3] + μ[1,5] ≥ 0
-μ[1,3] + μ[1,6] ≥ 0
-β_z[1] - μ[2,1] + μ[2,2] ≥ 0
-μ[2,1] + μ[2,3] ≥ 0
-β_z[2] - μ[2,2] + μ[2,4] ≥ 0
-μ[2,2] + μ[2,5] ≥ 0
-β_z[2] - μ[2,3] + μ[2,5] ≥ 0
-μ[2,3] + μ[2,6] ≥ 0
-λ + β_x[1] ≤ 0
-λ - β_x[1] ≤ 0
λ ≥ 0
(r[1] + λ * 0.0) - 1.0e-6 ≥ 0
((μ[1,7] - μ[1,4]) - (-1.0 * β_grp[1] - log(exp(r[1] + λ * 0.0) - 1.0)))
- 0.0 ≥ 0
(r[1] + λ * 0.5) - 1.0e-6 ≥ 0
((μ[1,7] - μ[1,4]) - -(log(exp(r[1] + λ * 0.5) - 1.0))) - 0.0 ≥ 0
(r[1] + λ * 1.0) - 1.0e-6 ≥ 0
((μ[1,7] - μ[1,5]) - (-1.0 * β_grp[1] - log(exp(r[1] + λ * 1.0) - 1.0)))
- 0.0 ≥ 0
(r[1] + λ * 1.5) - 1.0e-6 ≥ 0
((μ[1,7] - μ[1,5]) - -(log(exp(r[1] + λ * 1.5) - 1.0))) - 0.0 ≥ 0
(r[1] + λ * 2.0) - 1.0e-6 ≥ 0
((μ[1,7] - μ[1,6]) - (-1.0 * β_grp[1] - log(exp(r[1] + λ * 2.0) - 1.0)))
- 0.0 ≥ 0
(r[1] + λ * 2.5) - 1.0e-6 ≥ 0
((μ[1,7] - μ[1,6]) - -(log(exp(r[1] + λ * 2.5) - 1.0))) - 0.0 ≥ 0
(r[2] + λ * 2.5) - 1.0e-6 ≥ 0
((μ[2,7] - μ[2,4]) - (--1.0 * β_grp[1] - log(exp(r[2] + λ * 2.5) - 1.0)))
- 0.0 ≥ 0
(r[2] + λ * 2.0) - 1.0e-6 ≥ 0
((μ[2,7] - μ[2,4]) - -(log(exp(r[2] + λ * 2.0) - 1.0))) - 0.0 ≥ 0
(r[2] + λ * 1.5) - 1.0e-6 ≥ 0
((μ[2,7] - μ[2,5]) - (--1.0 * β_grp[1] - log(exp(r[2] + λ * 1.5) - 1.0)))
- 0.0 ≥ 0
(r[2] + λ * 1.0) - 1.0e-6 ≥ 0
((μ[2,7] - μ[2,5]) - -(log(exp(r[2] + λ * 1.0) - 1.0))) - 0.0 ≥ 0
(r[2] + λ * 0.5) - 1.0e-6 ≥ 0
((μ[2,7] - μ[2,6]) - (--1.0 * β_grp[1] - log(exp(r[2] + λ * 0.5) - 1.0)))
- 0.0 ≥ 0
(r[2] + λ * 0.0) - 1.0e-6 ≥ 0
((μ[2,7] - μ[2,6]) - -(log(exp(r[2] + λ * 0.0) - 1.0))) - 0.0 ≥ 0

```

In [31]: `optimize!(model)`

```

println("Termination status: ", termination_status(model))
println("Primal status: ", primal_status(model))

```

This is Ipopt version 3.14.19, running with linear solver MUMPS 5.8.1.

```

Number of nonzeros in equality constraint Jacobian...: 0
Number of nonzeros in inequality constraint Jacobian.: 119
Number of nonzeros in Lagrangian Hessian.....: 36

Total number of variables.....: 22
    variables with only lower bounds: 1
    variables with lower and upper bounds: 0
    variables with only upper bounds: 0
Total number of equality constraints.....: 0
Total number of inequality constraints.....: 40
    inequality constraints with only lower bounds: 38
    inequality constraints with lower and upper bounds: 0
    inequality constraints with only upper bounds: 2

iter      objective      inf_pr      inf_du   lg(mu)  ||d||    lg(rg) alpha_du alpha_
pr  ls
  0  1.1000000e+00  0.00e+00  1.50e+00  -1.0  0.00e+00    -  0.00e+00  0.00e+
00
  1  1.0651615e+00  0.00e+00  2.94e-01  -1.0  4.35e-01  -4.0  7.71e-01  1.00e+
00f 1
  2  4.0938881e-01  1.55e-02  1.35e-01  -1.7  1.13e+00  -4.5  7.34e-01  1.00e+
00f 1
  3  3.0306360e-01  0.00e+00  1.32e-01  -1.7  2.21e+00    -  8.50e-01  1.00e+
00h 1
  4  4.2268735e-01  0.00e+00  2.47e-01  -1.7  3.95e+00  -5.0  8.50e-01  1.00e+
00h 1
  5  7.6514495e-01  0.00e+00  3.40e-01  -1.7  1.63e+01    -  6.27e-01  9.00e-
01h 1
  6  8.0582093e-01  0.00e+00  4.36e-01  -1.7  8.21e+00    -  1.00e+00  1.00e+
00f 1
  7  8.2313419e-01  0.00e+00  1.42e+00  -1.7  7.13e+00    -  1.00e+00  1.00e+
00h 1
  8  8.3929669e-01  0.00e+00  1.75e-03  -1.7  8.65e-01  -5.4  1.00e+00  1.00e+
00h 1
  9  2.0307814e-01  2.79e-01  1.15e-01  -3.8  3.73e+01    -  1.00e+00  8.08e-
01f 1
iter      objective      inf_pr      inf_du   lg(mu)  ||d||    lg(rg) alpha_du alpha_
pr  ls
 10  1.5454031e-01  1.50e-01  1.13e-01  -3.8  2.41e+00  -5.9  8.41e-01  5.17e-
01h 1
 11  1.3150971e-01  9.76e-02  2.37e-02  -3.8  9.38e-01  -6.4  8.52e-01  8.03e-
01h 1
 12  1.3316334e-01  0.00e+00  4.52e-03  -3.8  9.64e-02  -6.9  1.00e+00  1.00e+
00h 1
 13  1.3348251e-01  0.00e+00  3.02e-05  -3.8  4.44e-02  -7.3  1.00e+00  1.00e+
00h 1
 14  1.3117326e-01  2.24e-04  1.45e-04  -5.7  8.21e-02    -  9.87e-01  9.68e-
01h 1
 15  1.3117411e-01  0.00e+00  4.32e-07  -5.7  1.50e-03  -7.8  1.00e+00  1.00e+
00h 1
 16  1.3114650e-01  0.00e+00  2.54e-08  -8.6  9.18e-04  -8.3  1.00e+00  1.00e+
00h 1
 17  1.3114651e-01  0.00e+00  2.81e-12  -8.6  1.65e-03  -8.8  1.00e+00  1.00e+
00h 1

```

Number of Iterations.....: 17

(scaled)

(unscaled)

Objective.....	1.3114651484864803e-01	1.3114651484864803e-
01		01
Dual infeasibility.....	2.8143952447955705e-12	2.8143952447955705e-
12		12
Constraint violation....	0.0000000000000000e+00	0.0000000000000000e+
00		00
Variable bound violation:	0.0000000000000000e+00	0.0000000000000000e+
00		00
Complementarity.....	2.5071538951071849e-09	2.5071538951071849e-
09		09
Overall NLP error.....	2.5071538951071849e-09	2.5071538951071849e-
09		09

Number of objective function evaluations	= 18
Number of objective gradient evaluations	= 18
Number of equality constraint evaluations	= 0
Number of inequality constraint evaluations	= 18
Number of equality constraint Jacobian evaluations	= 0
Number of inequality constraint Jacobian evaluations	= 18
Number of Lagrangian Hessian evaluations	= 17
Total seconds in IPOPT	= 0.010

EXIT: Optimal Solution Found.  
Termination status: LOCALLY\_SOLVED  
Primal status: FEASIBLE\_POINT

## larger example

```
In [ ]: using Random
using JuMP
using Ipopt

# -----
# If ZGEncodingInfo is already defined in your code, comment this out.
# -----
struct ZGEncodingInfo
    k_z::Vector{Int}      # number of categories for each categorical feat
    z_start::Vector{Int}  # starting index of each block in β_z (1-based)
    num_g::Int            # number of groups
end

# Sigmoid
σ(t) = 1 / (1 + exp(-t))

"""
    linpred(β₀, βₓ, βᵣ, βgrp, x, z, g, encinfo)

Compute β₀ + βₓᵀ x + βᵣᵀ φ_z(z) + βgrpᵀ φ_g(g)
using the SAME reduced coding as in the DRO model.
"""
function linpred(
    β₀::Real,
    βₓ::AbstractVector,
    βᵣ::AbstractVector,
    βgrp::AbstractVector,
    x::AbstractVector,
    z::AbstractVector,
    g::Int,
```

```

    encinfo::ZGEncodingInfo,
)
m      = length(encinfo.k_z)
num_g = encinfo.num_g

v = β₀ + dot(βx, x)

# categorical part
for ℓ in 1:m
    kℓ   = encinfo.k_z[ℓ]
    start = encinfo.z_start[ℓ]           # 1-based index in βz
    c     = z[ℓ]                         # category ∈ {1,...,kℓ}

    if c < kℓ
        idx = start + (c - 1)           # reduced dummy: last level is re
        v += βz[idx]
    end
end

# group part: reduced dummy, last group as reference
if g < num_g
    v += βgrp[g]
end

return v
end
"""
generate_synthetic_instance(; N=300)

```

Generate a synthetic dataset for testing the DRO LR model.

Returns:

X, Z, group, y, encinfo, β\_true

β\_true is a NamedTuple:  
 $(\beta_0 = \dots, \beta_x = \dots, \beta_z = \dots, \beta_{grp} = \dots)$

```

function generate_synthetic_instance(; N::Int = 500)
    Random.seed!(2025)

    # --- dimensions ---
    n_x   = 2                      # two numerical features
    k_z   = [3, 2]                  # two categorical features: sizes 3 and 2
    m     = length(k_z)
    num_g = 3                      # three groups

    p_z = sum(k_z .- 1)            # total length of βz (reduced dummy)
    p_g = num_g - 1                # length of βgrp

    # --- build encoding info for z ---
    z_start = zeros(Int, m)
    offset = 1
    for ℓ in 1:m
        z_start[ℓ] = offset
        offset += k_z[ℓ] - 1
    end
    encinfo = ZGEncodingInfo(k_z, z_start, num_g)

    # --- true parameters β* ---

```

```

β₀_true = -0.3
βx_true = [1.0, -0.7]
βz_true = [0.8, -0.5, 0.6] # length p_z = (3-1)+(2-1) = 3
βgrp_true = [0.5, -0.4] # length p_g = 2

@assert length(βx_true) == n_x
@assert length(βz_true) == p_z
@assert length(βgrp_true) == p_g

# --- sample features and labels ---
X = randn(N, n_x)
Z = zeros(Int, N, m)
group = zeros(Int, N)
y = zeros(Int, N)

for i in 1:N
    # categorical features
    for ℓ in 1:m
        Z[i, ℓ] = rand(1:k_z[ℓ])
    end

    # group index
    group[i] = rand(1:num_g)

    # linear predictor and label
    η = linpred(β₀_true, βx_true, βz_true, βgrp_true,
                view(X, i, :), view(Z, i, :), group[i], encinfo)
    p = σ(η)
    y[i] = rand() < p ? 1 : -1
end

β_true = (β₀ = β₀_true,
           βx = βx_true,
           βz = βz_true,
           βgrp = βgrp_true)

return X, Z, group, y, encinfo, β_true
end

-----
run_synthetic_experiment()

1. Generate synthetic data from a known β*.
2. Build the DRO LR graph-based model.
3. Solve it.
4. Print true vs estimated parameters.

Assumes you already defined `build_group_dro_graph_model`.
Also assumes that function returns `meta` containing JuMP variables
`β₀`, `β_x`, `β_z`, `β_grp`.
-----
function run_synthetic_experiment()
    # 1. data
    X, Z, group, y, encinfo, β_true = generate_synthetic_instance(N = 300
N, n_x = size(X)
m = size(Z, 2)
num_g = encinfo.num_g

    # 2. metric parameters (simple choice)
    delta = 1e-2 * ones(m) # δ_ℓ

```

```

A_group = ones(num_g)                      # A_g
B_group = ones(num_g)                      # B_g
C_group = ones(num_g)                      # C_g
gamma_x = 1e-2 * ones(n_x)                 # γ_j
ε           = 1e-4

# 3. build DRO model
optimizer = optimizer_with_attributes(Ipopt.Optimizer,
                                         "print_level" => 5)

model, meta = build_group_dro_graph_model(
    X, Z, group, y,
    encinfo,
    delta,
    A_group,
    B_group,
    C_group,
    gamma_x,
    ε,
    optimizer,
)

println("Model built. Start optimization...")
optimize!(model)
println("Termination status: ", termination_status(model))
println("Objective value:     ", objective_value(model))

# 4. extract parameters (adjust if your meta uses other field names)
β₀_hat   = value(meta.β₀)
βx_hat   = value.(meta.β_x)
βz_hat   = value.(meta.β_z)
βgrp_hat = value.(meta.β_grp)

println("\n==== True vs estimated parameters ===")
println("β₀  true = ", β_true.β₀, "  hat = ", β₀_hat)
println("βx  true = ", β_true.β_x, "  hat = ", βx_hat)
println("βz  true = ", β_true.β_z, "  hat = ", βz_hat)
println("βgrp true = ", β_true.βgrp, "  hat = ", βgrp_hat)

return model, meta, β_true, (β₀_hat, βx_hat, βz_hat, βgrp_hat)
end

# -----
# Example usage from REPL / notebook:
#
# include("your_dro_graph_code.jl")  # defines build_group_dro_graph_mod
# include("this_synthetic_test.jl")  # this file
model, meta, β_true, β_hat = run_synthetic_experiment()
# -----

```

Model built. Start optimization...  
 This is Ipopt version 3.14.19, running with linear solver MUMPS 5.8.1.

Number of nonzeros in equality constraint Jacobian....: 0  
 Number of nonzeros in inequality constraint Jacobian.: 24908  
 Number of nonzeros in Lagrangian Hessian.....: 8100

Total number of variables.....: 2409  
 variables with only lower bounds: 1  
 variables with lower and upper bounds: 0  
 variables with only upper bounds: 0  
 Total number of equality constraints.....: 0  
 Total number of inequality constraints.....: 7804  
 inequality constraints with only lower bounds: 7800  
 inequality constraints with lower and upper bounds: 0  
 inequality constraints with only upper bounds: 4

iter	objective	inf_pr	inf_du	lg(mu)	d	lg(rg)	alpha_du	alpha_ls
0	1.0001000e+00	0.00e+00	1.78e+00	-1.0	0.00e+00	-	0.00e+00	0.00e+00
00	0							
1	1.0305435e+00	0.00e+00	9.75e+00	-1.0	4.20e-01	-4.0	7.92e-01	1.44e-01
01f	1							
2	1.7004176e+00	0.00e+00	2.39e+01	-1.0	1.12e+00	-4.5	5.78e-01	1.00e+00
00f	1							
3	3.3735854e+00	0.00e+00	7.52e+00	-1.0	2.12e+00	-5.0	4.88e-01	1.00e+00
00f	1							
4	6.9104103e+00	0.00e+00	4.91e+00	-1.0	4.17e+00	-5.4	6.48e-01	1.00e+00
00f	1							
5	1.6512566e+01	0.00e+00	3.96e+00	-1.0	1.17e+01	-5.9	5.38e-01	1.00e+00
00f	1							
6	3.6127969e+01	0.00e+00	2.58e+00	-1.0	2.48e+01	-6.4	6.45e-01	1.00e+00
00f	1							
7	8.4408130e+01	0.00e+00	1.32e+00	-1.0	6.67e+01	-6.9	5.36e-01	1.00e+00
00f	1							
8	1.6307424e+02	0.00e+00	9.59e-02	-1.0	1.33e+02	-7.3	7.86e-01	1.00e+00
00f	1							
9	1.1217965e+02	0.00e+00	8.61e-01	-1.7	1.44e+02	-7.8	3.08e-01	1.00e+00
00f	1							
iter	objective	inf_pr	inf_du	lg(mu)	d	lg(rg)	alpha_du	alpha_ls
10	8.1734959e+01	0.00e+00	1.42e-02	-1.7	2.27e+02	-8.3	7.64e-01	1.00e+00
00f	1							
11	5.8149824e+01	0.00e+00	2.50e-02	-2.5	5.70e+01	-5.2	1.00e+00	7.69e-01
01f	1							
12	3.6778207e+01	0.00e+00	7.31e-02	-2.5	5.36e+01	-5.6	1.00e+00	5.27e-01
01f	1							
13	3.1838987e+01	0.00e+00	5.47e-02	-2.5	4.97e+01	-6.1	1.00e+00	2.93e-01
01f	1							
14	2.6793668e+01	0.00e+00	1.76e-01	-2.5	3.64e+01	-6.6	1.00e+00	9.67e-01
01h	1							
15	1.8768376e+01	0.00e+00	1.63e-01	-2.5	5.40e+01	-5.3	9.92e-01	6.80e-01
01f	1							
16	1.8845623e+01	0.00e+00	1.65e-02	-2.5	4.41e+00	-3.9	9.56e-01	1.00e+00
00h	1							
17	1.8830079e+01	0.00e+00	2.48e-02	-2.5	3.64e+00	-3.5	1.00e+00	1.00e+00
00h	1							
18	1.8668280e+01	0.00e+00	8.36e-03	-2.5	1.03e+00	-4.0	9.75e-01	1.00e+00
00h	1							
19	1.8106630e+01	0.00e+00	7.92e-03	-2.5	3.70e+00	-4.5	1.00e+00	1.00e+00

	iter	objective	inf_pr	inf_du	lg(mu)	d	lg(rg)	alpha_du	alpha_ls
00h	1								
00h	20	1.7883107e+01	0.00e+00	1.41e-02	-2.5	1.51e+00	-4.0	1.00e+00	1.00e+
00h	21	1.4572626e+01	0.00e+00	1.58e-01	-3.8	1.41e+01	-4.5	1.00e+00	6.79e-
01f	22	1.3626785e+01	0.00e+00	1.48e-01	-3.8	3.96e+00	-4.1	1.00e+00	7.05e-
01h	23	1.3222876e+01	0.00e+00	5.60e-02	-3.8	1.36e+00	-3.7	1.00e+00	1.00e+
00h	24	1.2780681e+01	0.00e+00	9.64e-02	-3.8	5.50e+00	-4.1	1.00e+00	3.97e-
01h	25	1.2531740e+01	0.00e+00	3.82e-02	-3.8	1.23e+00	-3.7	1.00e+00	1.00e+
00h	26	1.1807275e+01	0.00e+00	6.15e-02	-3.8	3.58e+00	-4.2	1.00e+00	1.00e+
00h	27	1.1590172e+01	0.00e+00	2.30e-02	-3.8	9.82e-01	-3.8	1.00e+00	1.00e+
00h	28	1.0864371e+01	2.37e-01	5.13e-02	-3.8	3.70e+00	-4.2	1.00e+00	1.00e+
00h	29	1.0654291e+01	0.00e+00	2.01e-02	-3.8	1.04e+00	-3.8	1.00e+00	1.00e+
00h	30	1.0261122e+01	7.62e+00	2.80e+00	-3.8	3.47e+00	-4.3	6.84e-01	5.66e-
01h	31	1.0225339e+01	6.65e+00	2.63e+00	-3.8	2.00e+00	-3.9	1.00e+00	1.77e-
01h	32	9.8802066e+00	3.88e+00	1.52e+00	-3.8	3.88e+00	-4.3	2.69e-01	5.10e-
01h	33	9.7281840e+00	1.90e+00	1.72e+00	-3.8	1.19e+00	-3.9	1.00e+00	7.71e-
01h	34	9.3355361e+00	6.46e-01	6.74e-01	-3.8	3.63e+00	-4.4	8.53e-01	6.44e-
01h	35	9.1683561e+00	1.05e-01	2.91e-01	-3.8	1.19e+00	-4.0	1.00e+00	1.00e+
00h	36	8.5759191e+00	5.61e-02	1.42e-01	-3.8	7.49e+00	-4.4	1.00e+00	8.80e-
01h	37	8.3955065e+00	8.69e-02	4.17e-02	-3.8	1.51e+00	-4.0	8.30e-01	9.05e-
00h	38	8.3523064e+00	0.00e+00	2.34e-02	-3.8	1.28e+00	-3.6	5.67e-01	1.00e+
00h	39	8.1671337e+00	0.00e+00	2.51e-02	-3.8	1.59e+00	-4.1	9.17e-01	1.00e+
00h	40	8.1085529e+00	0.00e+00	9.05e-03	-3.8	5.55e-01	-3.6	9.09e-01	1.00e+
00h	41	7.9086498e+00	0.00e+00	2.97e-02	-3.8	1.83e+00	-4.1	1.00e+00	1.00e+
00h	42	7.8451522e+00	0.00e+00	1.06e-02	-3.8	6.26e-01	-3.7	1.00e+00	1.00e+
00h	43	7.6243309e+00	0.00e+00	3.48e-02	-3.8	2.11e+00	-4.2	1.00e+00	1.00e+
00h	44	7.5548220e+00	0.00e+00	1.23e-02	-3.8	7.05e-01	-3.7	1.00e+00	1.00e+
00h	45	7.3057491e+00	0.00e+00	4.07e-02	-3.8	2.42e+00	-4.2	1.00e+00	1.00e+
00h	46	7.2293064e+00	0.00e+00	1.42e-02	-3.8	7.88e-01	-3.8	1.00e+00	1.00e+

	00h	1	47	6.9465335e+00	0.00e+00	4.70e-02	-3.8	2.76e+00	-4.3	1.00e+00	1.00e+
00h	1	48	6.8629700e+00	0.00e+00	1.60e-02	-3.8	8.73e-01	-3.8	1.00e+00	1.00e+	
00h	1	49	6.5437913e+00	3.56e-01	5.30e-02	-3.8	3.14e+00	-4.3	1.00e+00	1.00e+	
00h	1	iter	objective	inf_pr	inf_du	lg(mu)	d	lg(rg)	alpha_du	alpha_	
pr	ls	50	6.4545232e+00	0.00e+00	1.72e-02	-3.8	9.60e-01	-3.9	1.00e+00	1.00e+	
00h	1	51	6.1887297e+00	7.60e+00	2.56e+00	-3.8	3.49e+00	-4.4	9.65e-01	7.64e-	
01h	1	52	6.1827010e+00	7.31e+00	2.45e+00	-3.8	2.73e+00	-3.9	1.00e+00	6.32e-	
02h	1	53	5.9255068e+00	7.70e+00	5.57e+00	-3.8	4.72e+00	-4.4	2.50e-01	5.81e-	
01h	1	54	5.8968827e+00	6.35e+00	5.15e+00	-3.8	1.26e+00	-4.0	1.00e+00	2.76e-	
01h	1	55	5.6487756e+00	3.58e+00	3.42e+00	-3.8	6.11e+00	-4.5	3.22e-01	5.03e-	
01h	1	56	5.5610480e+00	1.85e+00	3.15e+00	-3.8	1.49e+00	-4.1	5.84e-01	8.12e-	
01h	1	57	5.2685068e+00	5.89e-01	1.49e+00	-3.8	5.44e+00	-4.5	8.15e-01	6.98e-	
01h	1	58	5.1778464e+00	3.10e-02	4.38e-01	-3.8	3.97e+00	-4.1	2.63e-01	1.00e+	
00h	1	59	4.8141398e+00	5.40e-02	3.22e-01	-3.8	5.19e+00	-4.6	7.95e-01	1.00e+	
00h	1	iter	objective	inf_pr	inf_du	lg(mu)	d	lg(rg)	alpha_du	alpha_	
pr	ls	60	4.8188893e+00	0.00e+00	1.87e-01	-3.8	8.22e-01	-2.3	1.00e+00	1.00e+	
00h	1	61	4.8196510e+00	0.00e+00	1.03e-01	-3.8	6.04e-01	-2.8	6.26e-01	1.00e+	
00h	1	62	4.8196283e+00	0.00e+00	7.37e-02	-3.8	8.24e-01	-2.4	5.91e-01	1.00e+	
00h	1	63	4.8197070e+00	0.00e+00	3.46e-02	-3.8	2.90e-01	-2.0	1.00e+00	1.00e+	
00h	1	64	4.8200865e+00	0.00e+00	4.08e-02	-3.8	2.34e+00	-2.5	2.51e-01	1.00e+	
00h	1	65	4.8210194e+00	0.00e+00	1.95e-02	-3.8	1.08e+00	-2.0	1.00e+00	1.00e+	
00h	1	66	4.8194361e+00	0.00e+00	1.07e-03	-3.8	1.40e-01	-2.5	1.00e+00	1.00e+	
00h	1	67	4.8135455e+00	0.00e+00	5.42e-04	-3.8	1.91e-01	-3.0	1.00e+00	1.00e+	
00h	1	68	4.7950527e+00	0.00e+00	2.13e-03	-3.8	6.12e-01	-3.5	1.00e+00	1.00e+	
00h	1	69	4.7415022e+00	0.00e+00	6.11e-03	-3.8	8.42e-01	-3.9	1.00e+00	1.00e+	
00h	1	iter	objective	inf_pr	inf_du	lg(mu)	d	lg(rg)	alpha_du	alpha_	
pr	ls	70	4.5646365e+00	0.00e+00	2.04e-02	-3.8	2.65e+00	-4.4	1.00e+00	1.00e+	
00h	1	71	4.5118938e+00	0.00e+00	7.11e-03	-3.8	8.77e-01	-4.0	1.00e+00	1.00e+	
00h	1	72	4.3263790e+00	0.00e+00	2.10e-02	-3.8	2.77e+00	-4.5	1.00e+00	1.00e+	
00h	1	73	4.2713646e+00	0.00e+00	7.30e-03	-3.8	9.04e-01	-4.0	1.00e+00	1.00e+	

```

00h 1
    74 4.0758659e+00 0.00e+00 2.17e-02 -3.8 2.91e+00 -4.5 1.00e+00 1.00e+
00h 1
    75 4.0181704e+00 0.00e+00 7.51e-03 -3.8 9.32e-01 -4.1 1.00e+00 1.00e+
00h 1
    76 3.8928050e+00 0.00e+00 1.67e-02 -3.8 3.15e+00 -4.6 1.00e+00 6.02e-
01h 1
    77 3.8290007e+00 0.00e+00 7.52e-03 -3.8 9.90e-01 -4.1 8.24e-01 1.00e+
00f 1
    78 3.7606327e+00 0.00e+00 1.27e-02 -3.8 3.93e+00 -4.6 1.00e+00 2.93e-
01h 1
    79 3.6895109e+00 0.00e+00 7.82e-03 -3.8 1.08e+00 -4.2 6.61e-01 1.00e+
00f 1
iter objective inf_pr inf_du lg(mu) ||d|| lg(rg) alpha_du alpha_
pr ls
    80 3.6524310e+00 0.00e+00 1.09e-02 -3.8 5.28e+00 -4.7 1.00e+00 1.37e-
01h 1
    81 3.5726715e+00 0.00e+00 8.44e-03 -3.8 1.18e+00 -4.2 5.85e-01 1.00e+
00f 1
    82 3.5514297e+00 0.00e+00 1.04e-02 -3.8 7.56e+00 -4.7 1.00e+00 6.64e-
02h 1
    83 3.4615955e+00 0.00e+00 9.35e-03 -3.8 1.31e+00 -4.3 5.53e-01 1.00e+
00f 1
    84 3.4494047e+00 0.00e+00 1.06e-02 -3.8 1.20e+01 -4.8 9.05e-01 3.10e-
02h 1
    85 3.3477148e+00 0.00e+00 1.05e-02 -3.8 1.47e+00 -4.3 5.46e-01 1.00e+
00f 1
    86 3.3408134e+00 0.00e+00 1.15e-02 -3.8 2.40e+01 -4.8 3.16e-01 1.29e-
02h 1
    87 3.2221593e+00 0.00e+00 1.17e-02 -3.8 1.84e+00 -4.4 1.00e+00 1.00e+
00f 1
    88 3.1837293e+00 0.00e+00 1.63e-02 -3.8 2.71e+01 -4.9 3.59e-01 4.92e-
02h 1
    89 3.0522335e+00 0.00e+00 1.52e-02 -3.8 1.90e+00 -4.4 9.69e-01 1.00e+
00h 1
iter objective inf_pr inf_du lg(mu) ||d|| lg(rg) alpha_du alpha_
pr ls
    90 3.0043228e+00 0.00e+00 3.75e-03 -3.8 9.68e-01 -4.0 1.00e+00 1.00e+
00h 1
    91 2.8609905e+00 0.00e+00 1.30e-02 -3.8 2.09e+00 -4.5 1.00e+00 1.00e+
00h 1
    92 2.8137790e+00 0.00e+00 4.71e-03 -3.8 7.30e-01 -4.1 1.00e+00 1.00e+
00h 1
    93 2.6515251e+00 0.00e+00 1.50e-02 -3.8 2.46e+00 -4.5 1.00e+00 1.00e+
00h 1
    94 2.5983388e+00 0.00e+00 5.38e-03 -3.8 7.87e-01 -4.1 1.00e+00 1.00e+
00h 1
    95 2.4125599e+00 0.00e+00 1.73e-02 -3.8 3.02e+00 -4.6 1.00e+00 1.00e+
00h 1
    96 2.3524139e+00 0.00e+00 6.21e-03 -3.8 8.84e-01 -4.2 1.00e+00 1.00e+
00h 1
    97 2.1374244e+00 0.00e+00 2.04e-02 -3.8 3.73e+00 -4.6 1.00e+00 1.00e+
00h 1
    98 2.0688179e+00 0.00e+00 7.26e-03 -3.8 1.04e+00 -4.2 1.00e+00 1.00e+
00h 1
    99 1.9155207e+00 0.00e+00 1.77e-02 -3.8 4.55e+00 -4.7 1.00e+00 6.10e-
01h 1
iter objective inf_pr inf_du lg(mu) ||d|| lg(rg) alpha_du alpha_
pr ls
    100 1.8389588e+00 0.00e+00 7.97e-03 -3.8 1.11e+00 -4.3 5.23e-01 1.00e+

```

	00h	1	101	1.7768936e+00	0.00e+00	1.24e-02	-3.8	5.47e+00	-4.8	1.00e+00	2.12e-
01h	1		102	1.6911970e+00	0.00e+00	8.11e-03	-3.8	1.25e+00	-4.3	3.95e-01	1.00e+
00f	1		103	1.5414948e+00	2.17e+00	2.38e-02	-3.8	5.48e+00	-4.8	1.00e+00	4.50e-
01h	1		104	1.4506536e+00	4.67e-01	1.46e-02	-3.8	1.42e+00	-4.4	2.42e-01	1.00e+
00h	1		105	1.2153165e+00	4.01e-01	2.04e-02	-3.8	6.69e+00	-4.9	5.84e-01	6.85e-
01h	1		106	1.1295937e+00	0.00e+00	7.50e-03	-3.8	3.68e+00	-4.4	3.09e-01	1.00e+
00h	1		107	1.0086186e+00	5.64e-02	1.20e-02	-3.8	5.54e+00	-4.9	6.33e-01	4.52e-
01h	1		108	8.3160330e-01	5.27e-01	2.17e-02	-3.8	1.81e+01	-5.4	3.27e-01	2.59e-
01h	1		109	7.7494902e-01	1.19e-01	6.44e-03	-3.8	2.92e+00	-5.0	1.00e+00	1.00e+
00h	1		iter	objective	inf_pr	inf_du	lg(mu)	d	lg(rg)	alpha_du	alpha_
pr	ls		110	7.6980750e-01	0.00e+00	8.82e-04	-3.8	2.26e+00	-5.4	1.00e+00	1.00e+
00h	1		111	7.7147083e-01	0.00e+00	6.24e-05	-3.8	2.05e+00	-5.9	1.00e+00	1.00e+
00h	1		112	7.7124884e-01	0.00e+00	4.88e-04	-3.8	1.97e+00	-6.4	1.00e+00	1.00e+
00h	1		113	6.0508542e-01	7.64e-02	5.92e-04	-5.7	2.93e+00	-6.9	7.85e-01	7.87e-
01f	1		114	5.7902958e-01	1.89e-03	1.22e-04	-5.7	7.57e-01	-7.3	9.88e-01	1.00e+
00h	1		115	5.7806101e-01	0.00e+00	1.13e-04	-5.7	3.42e+00	-7.8	1.00e+00	1.00e+
00h	1		116	5.7789302e-01	0.00e+00	1.52e-04	-5.7	1.56e+00	-7.4	1.00e+00	1.00e+
00h	1		117	5.7786996e-01	0.00e+00	2.76e-05	-5.7	3.03e-01	-7.0	1.00e+00	1.00e+
00h	1		118	5.7777080e-01	0.00e+00	6.11e-05	-5.7	1.02e+00	-7.4	1.00e+00	1.00e+
00h	1		119	5.7774783e-01	0.00e+00	1.43e-05	-5.7	2.82e-01	-7.0	1.00e+00	1.00e+
00h	1		iter	objective	inf_pr	inf_du	lg(mu)	d	lg(rg)	alpha_du	alpha_
pr	ls		120	5.7540466e-01	1.64e-03	5.67e-05	-8.6	1.44e+00	-7.5	9.47e-01	9.97e-
01h	1		121	5.7535308e-01	0.00e+00	2.29e-05	-8.6	4.63e-01	-7.1	8.84e-01	1.00e+
00h	1		122	5.7518739e-01	0.00e+00	4.85e-05	-8.6	1.69e+00	-7.5	1.00e+00	1.00e+
00h	1		123	5.7483182e-01	0.00e+00	1.10e-04	-8.6	1.20e+01	-8.0	3.34e-01	3.02e-
01h	1		124	5.7464428e-01	0.00e+00	1.06e-04	-8.6	1.73e+01	-8.5	4.13e-01	1.11e-
01h	1		125	5.7407040e-01	0.00e+00	8.44e-05	-8.6	2.78e+01	-9.0	4.59e-01	2.11e-
01f	1		126	5.7238571e-01	0.00e+00	6.61e-05	-8.6	7.92e+01	-9.5	7.72e-01	2.17e-
01h	1		127	5.4992563e-01	0.00e+00	2.24e-06	-8.6	2.37e+02	-9.9	1.00e+00	9.66e-
01h	1		128	5.4926259e-01	2.55e-02	2.22e-06	-8.6	7.08e+02	-10.4	1.00e+00	9.56e-

```

03h 1
 129 5.4855453e-01 3.34e-02 1.70e-03 -8.6 7.55e+00 -10.9 8.29e-01 1.00e+
00h 1
iter   objective     inf_pr     inf_du   lg(mu) ||d||   lg(rg) alpha_du alpha_
pr  ls
 130 5.4870612e-01 1.65e-02 1.31e-03 -8.6 1.98e+00 -11.4 2.34e-01 5.10e-
01h 1
 131 5.4879600e-01 7.18e-03 2.24e-04 -8.6 9.56e-01 -11.8 8.26e-01 5.64e-
01h 1
 132 5.4886485e-01 5.26e-04 4.87e-05 -8.6 2.34e+00 -11.4 7.85e-01 9.24e-
01h 1
 133 5.4887077e-01 0.00e+00 3.74e-10 -8.6 1.26e+00 -11.9 1.00e+00 9.99e-
01h 1

```

Number of Iterations....: 133

	(scaled)	(unscaled)
Objective.....: 5.4887076556059111e-01	5.4887076556059111e-01	5.4887076556059111e-01
Dual infeasibility.....: 3.7365899622253190e-10	3.7365899622253190e-10	3.7365899622253190e-10
Constraint violation....: 0.0000000000000000e+00	0.0000000000000000e+00	0.0000000000000000e+00
Variable bound violation: 0.0000000000000000e+00	0.0000000000000000e+00	0.0000000000000000e+00
Complementarity.....: 3.2711975698061975e-09	3.2711975698061975e-09	3.2711975698061975e-09
Overall NLP error.....: 3.2711975698061975e-09	3.2711975698061975e-09	3.2711975698061975e-09

Number of objective function evaluations	= 134
Number of objective gradient evaluations	= 134
Number of equality constraint evaluations	= 0
Number of inequality constraint evaluations	= 134
Number of equality constraint Jacobian evaluations	= 0
Number of inequality constraint Jacobian evaluations	= 134
Number of Lagrangian Hessian evaluations	= 133
Total seconds in IPOPT	= 6.498

EXIT: Optimal Solution Found.

Termination status: LOCALLY\_SOLVED

Objective value: 0.5488707655605911

==== True vs estimated parameters ===

```

β₀  true = -0.3  hat = -0.26966141856599213
βₓ  true = [1.0, -0.7]  hat = [0.7268207209228954, -0.8371227622608812]
βₜ  true = [0.8, -0.5, 0.6]  hat = [0.7977649681015079, -0.1533352980175
6676, 0.5879875082314706]
βgrp true = [0.5, -0.4]  hat = [0.8779178610381262, -0.6205919690109933]

```

```
(A JuMP Model
|- solver: Ipopt
|- objective_sense: MIN_SENSE
  |- objective_function_type: AffExpr
|- num_variables: 2409
|- num_constraints: 7805
  |- AffExpr in MOI.GreaterThan{Float64}: 2400
  |- AffExpr in MOI.LessThan{Float64}: 4
  |- VariableRef in MOI.GreaterThan{Float64}: 1
  |- Nonlinear: 5400
|- Names registered in the model
  |- :r, :β₀, :β.grp, :β_x, :β_z, :λ, :μ, (encinfo = ZGEncodingInfo([3, 2], [1, 3], 3), dags = SampleDAG[SampleDAG(1, [(0, 0.0), (1, 0.01), (1, 0.0), (2, 0.01), (2, 0.02), (2, 0.0), (3, 0.0)], Arc[Arc(1, 2, CatArc(1, 1, 0.0, 0.01)), Arc(1, 2, CatArc(1, 2, 0.0, 0.01)), Arc(1, 3, CatArc(1, 3, 0.0, 0.0)), Arc(2, 4, CatArc(2, 1, 0.01, 0.01)), Arc(2, 5, CatArc(2, 2, 0.01, 0.02)), Arc(3, 6, CatArc(2, 1, 0.0, 0.0)), Arc(3, 4, CatArc(2, 2, 0.0, 0.01)), Arc(4, 7, TermArc(0.01, 1)), Arc(4, 7, TermArc(0.01, 2)), Arc(4, 7, TermArc(0.01, 3)), Arc(5, 7, TermArc(0.02, 1)), Arc(5, 7, TermArc(0.02, 2)), Arc(5, 7, TermArc(0.02, 3)), Arc(6, 7, TermArc(0.0, 1)), Arc(6, 7, TermArc(0.0, 2)), Arc(6, 7, TermArc(0.0, 3))], 1, 7), SampleDAG(2, [(0, 0.0), (1, 0.01), (1, 0.0), (2, 0.01), (2, 0.02), (2, 0.0), (3, 0.0)], Arc[Arc(1, 2, CatArc(1, 1, 0.0, 0.01)), Arc(1, 2, CatArc(1, 2, 0.0, 0.01)), Arc(1, 3, CatArc(1, 3, 0.0, 0.0)), Arc(2, 4, CatArc(2, 1, 0.01, 0.01)), Arc(2, 5, CatArc(2, 2, 0.01, 0.02)), Arc(3, 6, CatArc(2, 1, 0.0, 0.0)), Arc(3, 4, CatArc(2, 2, 0.0, 0.01)), Arc(4, 7, TermArc(0.01, 1)), Arc(4, 7, TermArc(0.01, 2)), Arc(4, 7, TermArc(0.01, 3)), Arc(5, 7, TermArc(0.02, 1)), Arc(5, 7, TermArc(0.02, 2)), Arc(5, 7, TermArc(0.02, 3)), Arc(6, 7, TermArc(0.0, 1)), Arc(6, 7, TermArc(0.0, 2)), Arc(6, 7, TermArc(0.0, 3))], 1, 7), SampleDAG(3, [(0, 0.0), (1, 0.01), (1, 0.0), (2, 0.01), (2, 0.02), (2, 0.0), (3, 0.0)], Arc[Arc(1, 2, CatArc(1, 1, 0.0, 0.01)), Arc(1, 2, CatArc(1, 2, 0.0, 0.01)), Arc(1, 3, CatArc(1, 3, 0.0, 0.0)), Arc(2, 4, CatArc(2, 1, 0.01, 0.01)), Arc(2, 5, CatArc(2, 2, 0.01, 0.02)), Arc(3, 6, CatArc(2, 1, 0.0, 0.0)), Arc(3, 4, CatArc(2, 2, 0.0, 0.01)), Arc(4, 7, TermArc(0.01, 1)), Arc(4, 7, TermArc(0.01, 2)), Arc(4, 7, TermArc(0.01, 3)), Arc(5, 7, TermArc(0.02, 1)), Arc(5, 7, TermArc(0.02, 2)), Arc(5, 7, TermArc(0.02, 3)), Arc(6, 7, TermArc(0.0, 1)), Arc(6, 7, TermArc(0.0, 2)), Arc(6, 7, TermArc(0.0, 3))], 1, 7), SampleDAG(4, [(0, 0.0), (1, 0.0), (1, 0.01), (2, 0.01), (2, 0.0), (2, 0.02), (3, 0.0)], Arc[Arc(1, 2, CatArc(1, 1, 0.0, 0.01)), Arc(1, 2, CatArc(1, 2, 0.0, 0.01)), Arc(1, 3, CatArc(1, 3, 0.0, 0.0)), Arc(2, 4, CatArc(2, 1, 0.01, 0.01)), Arc(2, 5, CatArc(2, 2, 0.01, 0.02)), Arc(3, 6, CatArc(2, 1, 0.0, 0.0)), Arc(3, 4, CatArc(2, 2, 0.0, 0.01)), Arc(4, 7, TermArc(0.01, 1)), Arc(4, 7, TermArc(0.01, 2)), Arc(4, 7, TermArc(0.01, 3)), Arc(5, 7, TermArc(0.02, 1)), Arc(5, 7, TermArc(0.02, 2)), Arc(5, 7, TermArc(0.02, 3)), Arc(6, 7, TermArc(0.0, 1)), Arc(6, 7, TermArc(0.0, 2)), Arc(6, 7, TermArc(0.0, 3))], 1, 7), SampleDAG(5, [(0, 0.0), (1, 0.01), (1, 0.0), (2, 0.02), (2, 0.01), (2, 0.0), (3, 0.0)], Arc[Arc(1, 2, CatArc(1, 1, 0.0, 0.01)), Arc(1, 2, CatArc(1, 2, 0.0, 0.01)), Arc(1, 3, CatArc(1, 3, 0.0, 0.0)), Arc(2, 4, CatArc(2, 1, 0.01, 0.02)), Arc(2, 5, CatArc(2, 2, 0.01, 0.01)), Arc(3, 6, CatArc(2, 1, 0.0, 0.0)), Arc(4, 7, TermArc(0.02, 1)), Arc(4, 7, TermArc(0.02, 2)), Arc(4, 7, TermArc(0.02, 3)), Arc(5, 7, TermArc(0.0, 1)), Arc(5, 7, TermArc(0.0, 2)), Arc(5, 7, TermArc(0.0, 3)), Arc(6, 7, TermArc(0.02, 1)), Arc(6, 7, TermArc(0.02, 2)), Arc(6, 7, TermArc(0.02, 3))], 1, 7), SampleDAG(6, [(0, 0.0), (1, 0.01), (1, 0.0), (2, 0.02), (2, 0.01), (2, 0.0), (3, 0.0)], Arc[Arc(1, 2, CatArc(1, 1, 0.0, 0.01)), Arc(1, 3, CatArc(1, 2, 0.0, 0.01)), Arc(1, 2, CatArc(1, 3, 0.0, 0.01)), Arc(2, 4, CatArc(2, 1, 0.01, 0.02)), Arc(2, 5, CatArc(2, 2, 0.01, 0.01)), Arc(3, 6, CatArc(2, 1, 0.0, 0.0)), Arc(4, 7, TermArc(0.01, 1)), Arc(5, 7, TermArc(0.01, 1)), Arc(6, 7, TermArc(0.0, 1)), Arc(6, 7, TermArc(0.0, 2)), Arc(6, 7, TermArc(0.0, 3))], 1, 7)
```





```
ableRef[β_x[1], β_x[2]], β_z = VariableRef[β_z[1], β_z[2], β_z[3]], β_grp
= VariableRef[β_grp[1], β_grp[2]], λ = λ, r = VariableRef[r[1], r[2], r
[3], r[4], r[5], r[6], r[7], r[8], r[9], r[10] ... r[291], r[292], r[293],
r[294], r[295], r[296], r[297], r[298], r[299], r[300]], μ = VariableRef[μ
[1,1] μ[1,2] ... μ[1,6] μ[1,7]; μ[2,1] μ[2,2] ... μ[2,6] μ[2,7]; ... ; μ[299,1]
μ[299,2] ... μ[299,6] μ[299,7]; μ[300,1] μ[300,2] ... μ[300,6] μ[300,7]]], (β₀
= -0.3, βx = [1.0, -0.7], βz = [0.8, -0.5, 0.6], βgrp = [0.5, -0.4]), (-0.
26966141856599213, [0.7268207209228954, -0.8371227622608812], [0.797764968
1015079, -0.15333529801756676, 0.5879875082314706], [0.8779178610381262,
-0.6205919690109933]))
```

# Experiments

## some helper functions

In [79]:

```
#####
# 1. Parameter containers
#####
"""
DROParams

Stores all hyperparameters that define the subgroup
Wasserstein ground metric and DRO radius:

d(ξ^i, ξ) =
    A_g * Σ_j γ_j |x_j - x_j^i|
+ B_g * Σ_ℓ δ_ℓ 1[z_ℓ ≠ z_ℓ^i]
+ C_g * 1[g ≠ g^i]
+ ∞ * 1[y ≠ y^i].
"""

struct DROParams
    delta::Vector{Float64}          # length m, δ_ℓ for categorical features
    A_group::Vector{Float64}         # length num_g, A_g for continuous part
    B_group::Vector{Float64}         # length num_g, B_g for categorical part
    C_group::Vector{Float64}         # length num_g, C_g for group jumps
    gamma_x::Vector{Float64}         # length n_x, γ_j for continuous features
    epsilon::Float64                # Wasserstein radius ε
end

"""
LogitParams

Hyperparameters for standard (non-robust) logistic regression.

"""

struct LogitParams
    lambda_l2::Float64             # L2 regularization coefficient
end

"""
PerturbParams

Extra knobs for generating perturbed test sets.

The actual "cost" of moving features comes from DROParams
(via A_g * γ_j and B_g * δ_ℓ). Here we only store:

```

```

- mode : scenario label
- w    : half-width for continuous noise U(-w, w)
"""

struct PerturbParams
    mode::Symbol
    w::Float64
end

#####
# 2. Simple helper: DRO → Perturb
#####

"""

dro_to_perturb_params(dro; w=0.4, mode=:generic)

Create a PerturbParams object from a DROParams, choosing only
the non-metric perturbation hyperparameters (mode, w).

The metric itself (A_group, B_group, gamma_x, delta) is used
directly inside the perturbation routine, not stored here.
"""

function dro_to_perturb_params(
    dro::DROParams;
    w::Float64 = 0.4,
    mode::Symbol = :generic,
)
    return PerturbParams(mode, w)
end

```

`dro_to_perturb_params`

In [80]:

```

#####
# 3. Model fitting (DRO vs Logistic) + shared prediction
#####

using JuMP
using Ipopt

"""

fit_dro_model(X, Z, group, y, encinfo, params; optimizer)

Fit the subgroup-Wasserstein DRO logistic regression using
`build_group_dro_graph_model` and return coefficients for prediction.

```

**Inputs**

-----

- X : N×n<sub>x</sub> continuous features
- Z : N×m integer-coded categorical features
- group : length-N group indices in {1,...,num<sub>g</sub>}
- y : length-N labels in {-1,+1}
- encinfo : ZGEncodingInfo (reduced encoding info)
- params : DROParams
- optimizer : JuMP optimizer constructor (default Ipopt)

**Returns**

-----

NamedTuple (β₀, β<sub>x</sub>, β<sub>z</sub>, β<sub>grp</sub>, encinfo)

"""

**function** fit\_dro\_model(

```

X::AbstractMatrix,
Z::AbstractMatrix{<:Integer},
group::AbstractVector{<:Integer},
y::AbstractVector{<:Integer},
encinfo::ZGEncodingInfo,
params::DROParams;
optimizer = optimizer_with_attributes(Ipopt.Optimizer, "print_level"
)
model, meta = build_group_dro_graph_model(
    X,
    Z,
    group,
    y,
    encinfo,
    params.delta,
    params.A_group,
    params.B_group,
    params.C_group,
    params.gamma_x,
    params.epsilon,
    optimizer,
)
optimize!(model)

β̂     = value(meta.β₀)
β̂_x   = value.(meta.β_x)
β̂_z   = value.(meta.β_z)
β̂_grp = value.(meta.β_grp)

return (
    β₀      = β̂,
    β_x    = β̂_x,
    β_z    = β̂_z,
    β_grp = β̂_grp,
    encinfo = encinfo,
)
end

"""
fit_logistic_model(Xtr, Zenc_tr, Genc_tr, ytr, encinfo, params; optim
Fit standard logistic regression with L2 regularization, using the
same reduced encoding as the DRO model:


$$f_{\beta}(x, z, g) = \beta_0 + \beta_x^T x + \beta_z^T \phi_z(z) + \beta_{grp}^T \phi_g(g)$$

and


$$\min_{\beta} (1/N) \sum \log(1 + \exp(-y_i f_{\beta}(x_i, z_i, g_i))) + (\lambda/2)(\|\beta_x\|^2 + \|\beta_z\|^2 + \|\beta_{grp}\|^2),$$

with  $\lambda = \text{params}.lambda_l2.$ 
"""

function fit_logistic_model(
    Xtr::AbstractMatrix,
    Zenc_tr::AbstractMatrix,
    Genc_tr::AbstractMatrix,
    ytr::AbstractVector{<:Integer},
)

```

```

encinfo:::ZEncodingInfo,
params:::LogitParams;
optimizer = optimizer_with_attributes(Ipopt.Optimizer, "print_level"
)
N, n_x  = size(Xtr)
N2, p_z = size(Zenc_tr)
N3, p_g = size(Genc_tr)

@assert N2 == N "Zenc_tr must have same rows as Xtr"
@assert N3 == N "Genc_tr must have same rows as Xtr"

λ = params.lambda_l2

model = Model(optimizer)

@variable(model, β₀)
@variable(model, β_x[1:n_x])
@variable(model, β_z[1:p_z])
@variable(model, β_grp[1:p_g])

# Logistic loss + L2 penalty (no penalty on intercept)
@NLobjective(model, Min,
(1.0 / N) * sum(
    log(1 + exp(-ytr[i] * (
        β₀
        + sum(β_x[j] * Xtr[i, j] for j in 1:n_x)
        + sum(β_z[k] * Zenc_tr[i, k] for k in 1:p_z)
        + sum(β_grp[h] * Genc_tr[i, h] for h in 1:p_g)
    )))
    for i in 1:N
)
+ (λ / 2.0) * (
    sum(β_x[j]^2 for j in 1:n_x) +
    sum(β_z[k]^2 for k in 1:p_z) +
    sum(β_grp[h]^2 for h in 1:p_g)
)
)
)

optimize!(model)

return (
    β₀      = value(β₀),
    β_x     = value.(β_x),
    β_z     = value.(β_z),
    β_grp   = value.(β_grp),
    encinfo = encinfo,
)
end

#####
# Shared prediction functions
#####

# Logistic link
σ(t) = 1.0 / (1.0 + exp(-t))

.....
predict_scores(β, X, Z, group)

```

```

Compute linear scores  $f_\beta(x, z, g)$  for each row, using the
encoding described by  $\beta.\text{encinfo}$ .
"""
function predict_scores(
    β,
    X::AbstractMatrix,
    Z::AbstractMatrix{<:Integer},
    group::AbstractVector{<:Integer},
)
    N, n_x = size(X)
    k_z      = β.encinfo.k_z
    z_start  = β.encinfo.z_start
    num_g    = β.encinfo.num_g

    m = length(k_z)
    @assert size(Z, 2) == m
    @assert length(group) == N

    scores = zeros(Float64, N)

    for i in 1:N
        s = β.β₀ + dot(β.β_x, view(X, i, :))

        # categorical part
        for ℓ in 1:m
            val = Z[i, ℓ]
            k_ℓ = k_z[ℓ]
            if val < k_ℓ
                idx = z_start[ℓ] + (val - 1)
                s += β.β_z[idx]
            end
        end

        # group part
        g_i = group[i]
        if g_i < num_g
            s += β.β_grp[g_i]
        end

        scores[i] = s
    end

    return scores
end

"""
predict_proba(β, X, Z, group)

Return  $P(y=+1 | x, z, g) = \sigma(f_\beta(x, z, g))$  for each row.
"""
predict_proba(β, X, Z, group) = σ.(predict_scores(β, X, Z, group))

"""
predict_label(β, X, Z, group)

Return hard labels in {-1,+1} using  $\text{sign}(f_\beta)$ .
"""
function predict_label(
    β,
    X::AbstractMatrix,
)

```

```

Z::AbstractMatrix{<:Integer},
group::AbstractVector{<:Integer},
)
scores = predict_scores(β, X, Z, group)
yhat = Vector{Int}(undef, length(scores))
@inbounds for i in eachindex(scores)
    yhat[i] = scores[i] >= 0 ? 1 : -1
end
return yhat
end

```

`predict_label`

```

In [88]: #####
# 4. Group-aware test-set perturbation driven by metric
#####

using Random
using Distributions # for Laplace

#####
perturb_testset(X, Z, group, y, encinfo, dro, pert; rng)

Generate one perturbed test set ( $\tilde{X}$ ,  $\tilde{Z}$ ,  $\tilde{g}$ ,  $\tilde{y}$ ) from
( $X$ ,  $Z$ ,  $group$ ,  $y$ ), using the subgroup Wasserstein metric:

d( $\xi^i$ ,  $\xi$ ) =
A_g  $\sum_j \gamma_j |x_j - x_j^{i*}|$ 
+ B_g  $\sum_\ell \delta_\ell 1[z_\ell \neq z_\ell^{i*}]$ 
+ C_g  $1[g \neq g^i]$ .
```

We use these products  $A_g \cdot \gamma_j$ ,  $B_g \cdot \delta_\ell$ ,  $C_g$  to control the amount of noise:

Continuous features:

For sample  $i$  and feature  $j$  (group  $g_i$ ):

```

cost_x = A_{g_i} * γ_j
scale = w / cost_x
Δ_raw ~ Laplace(0, scale)
Δ = clamp(Δ_raw, -w, w)
X̃[i,j] = X[i,j] + Δ

```

Categorical features:

For sample  $i$  and feature  $\ell$  (group  $g_i$ ):

```

cost_z = B_{g_i} * δ_ℓ
q = 1 - exp(-w / cost_z) ∈ (0,1)
With prob 1-q: keep  $z_\ell^{i*}$ 
With prob q : change uniformly to one of the other levels.

```

Group index:

For sample  $i$ :

```

cost_g = C_{g_i}
qg = 1 - exp(-w / cost_g) ∈ (0,1)
With prob 1-qg: keep  $g_i$ 
With prob qg : change to a different group, uniformly.

```

Inputs:

- $X$ ,  $Z$ ,  $group$ ,  $y$  : test data
- $encinfo$  : ZGEncodingInfo ( $k_z$ ,  $num_g$ )
- $dro$  : DROParams ( $\delta$ ,  $A_group$ ,  $B_group$ ,  $C_group$ ,  $gamma_x$ ,

```

- pert : PerturbParams (mode, w)
- rng : random number generator

>Returns:
-  $\tilde{X}$ ,  $\tilde{Z}$ ,  $\tilde{g}$ ,  $\tilde{y}$  : perturbed copies (same shapes as inputs)
"""

function perturb_testset(
    X::AbstractMatrix,
    Z::AbstractMatrix{<:Integer},
    group::AbstractVector{<:Integer},
    y::AbstractVector,
    encinfo::ZGEncodingInfo,
    dro::DROPParams,
    pert::PerturbParams;
    rng = Random.default_rng(),
)
    N, n_x = size(X)
    _, m = size(Z)
    k_z = encinfo.k_z
    num_g = encinfo.num_g

    @assert length(dro.gamma_x) == n_x
    @assert length(dro.delta) == m
    @assert length(dro.A_group) == num_g
    @assert length(dro.B_group) == num_g
    @assert length(dro.C_group) == num_g

     $\tilde{X}$  = copy(X)
     $\tilde{Z}$  = copy(Z)
     $\tilde{g}$  = copy(group)
     $\tilde{y}$  = copy(y)

    w = pert.w
    @assert w > 0 "pert.w must be positive"

    for i in 1:N
        g_i = group[i]
        @assert 1 ≤ g_i ≤ num_g

        A_gi = dro.A_group[g_i]
        B_gi = dro.B_group[g_i]
        C_gi = max(dro.C_group[g_i], 1e-6) # avoid zero

        # -----
        # 1. Continuous features
        # -----
        for j in 1:n_x
            cost_x = max(A_gi * dro.gamma_x[j], 1e-6)
            scale = w / cost_x
            Δ_raw = rand(rng, Laplace(0.0, scale))
            # clip extreme jumps; still  $O(w / (A_gi))$ 
            Δ = clamp(Δ_raw, -w, w)
             $\tilde{X}$ [i, j] = X[i, j] + Δ
        end

        # -----
        # 2. Categorical features
        # -----
        for ℓ in 1:m
            k_ℓ = k_z[ℓ]
        end
    end
end

```

```

cost_z = max(B_gi * dro.delta[l], 1e-6)
# higher cost_z → smaller q
q = 1.0 - exp(-w / cost_z)
q = clamp(q, 0.0, 0.99)

if rand(rng) < q
    old = Z[i, l]
    # choose any other category uniformly
    choices = Vector{Int}(undef, k_l - 1)
    idx = 1
    for c in 1:k_l
        if c != old
            choices[idx] = c
            idx += 1
        end
    end
    Ź[i, l] = rand(rng, choices)
end

# -----
# 3. Group index
# -----
# If C_g very large, qg is very small (hard to move group)
qg = 1.0 - exp(-w / C_gi)
qg = clamp(qg, 0.0, 0.99)

if rand(rng) < qg
    choices_g = Vector{Int}(undef, num_g - 1)
    idxg = 1
    for gg in 1:num_g
        if gg != g_i
            choices_g[idxg] = gg
            idxg += 1
        end
    end
    ġ[i] = rand(rng, choices_g)
end

return Ÿ, Ź, ġ, Ÿ
end

```

perturb\_testset

In [89]:

```

#####
# 5. Evaluation metrics: AUC and ACE
#####

using Statistics

"""
    auc_binary(scores, y)

Binary AUC (ROC area).

- `scores`: model scores (logits or probabilities).
- `y`: true labels, positives are `y > 0`.

Implements the classic rank-based formula (Mann-Whitney U):

```

```

AUC = ( $\sum \text{rank}(\text{pos}) - P(P+1)/2) / (P \cdot N)$ ,

where  $P = \# \text{positives}$ ,  $N = \#\text{negatives}$ .
Returns 0.5 if all labels are the same.
"""

function auc_binary(scores::AbstractVector{<:Real},
                     y::AbstractVector{<:Real})
    @assert length(scores) == length(y)
    n = length(scores)

    # map labels to 0/1
    y01 = Vector{Int}(undef, n)
    for i in 1:n
        y01[i] = y[i] > 0 ? 1 : 0
    end

    P = sum(y01)
    N = n - P
    (P == 0 || N == 0) && return 0.5

    # ranks of scores (ascending)
    idx = sortperm(scores)
    ranks = similar(scores, Float64)
    for (r, i) in enumerate(idx)
        ranks[i] = r
    end

    # sum of ranks for positives
    sum_r_pos = 0.0
    for i in 1:n
        if y01[i] == 1
            sum_r_pos += ranks[i]
        end
    end

    auc = (sum_r_pos - P * (P + 1) / 2) / (P * N)
    return auc
end

"""

ace_binary(probs, y; B = 10)

Adaptive Calibration Error (ACE) for binary classification.

- `probs`: predicted  $P(y=1) \in [0,1]$ .
- `y`: true labels, positives are `y > 0`.
- `B`: number of equal-mass bins (default 10).

Procedure:
1. Sort by `probs`.
2. Split into B bins with (almost) equal size.
3. For each bin:
   | mean(probs) - mean(labels) |
4. ACE = average of these B values.

Lower ACE = better calibration.
"""

function ace_binary(probs::AbstractVector{<:Real},

```

```

y::AbstractVector{<:Real};
B::Int = 10
@assert length(probs) == length(y)
n = length(probs)
n == 0 && return 0.0

# map labels to 0/1
y01 = Vector{Float64}(undef, n)
for i in 1:n
    y01[i] = y[i] > 0 ? 1.0 : 0.0
end

B = min(B, n)

# sort by predicted probability
idx = sortperm(probs)

base = div(n, B)
extra = n % B

start = 1
err_sum = 0.0
for b in 1:B
    sz = base + (b <= extra ? 1 : 0)
    stop = start + sz - 1

    inds = idx[start:stop]
    mean_conf = mean(@view probs[inds])
    mean_acc = mean(@view y01[inds])

    err_sum += abs(mean_acc - mean_conf)
    start = stop + 1
end

return err_sum / B
end

#####
default_metrics(probs, y)

```

Convenience wrapper:

- `auc`: ROC AUC (higher is better)
- `ace`: Adaptive Calibration Error (lower is better)

```

#####
default_metrics(probs::AbstractVector{<:Real},
                y::AbstractVector{<:Real}) = (
    auc = auc_binary(probs, y),
    ace = ace_binary(probs, y; B = 10),
)

```

default\_metrics

```

In [90]: #####
# 6. High-level experiment driver (UNPAIRED)
#####

#####
run_experiment(

```

```

        X, Z, group, y;
        train_ratio = 0.7,
        method = :dro,
        dro_params = nothing,
        logit_params = nothing,
        pert_params::PerturbParams,
        n_splits = 1,
        n_pert_per_split = 100,
        metric_fun = default_metrics,
        rng = Random.default_rng(),
        optimizer = optimizer_with_attributes(Ipopt.Optimizer, "print_level"
    )
)

```

Run a robustness experiment for one model type:

- `method = :dro` → subgroup DRO logistic regression
- `method = :logistic` → standard L2 logistic regression

For each split  $s = 1, \dots, n_{\text{splits}}$ :

1. Randomly split data into train / test (ratio = `train\_ratio`).
2. Fit the chosen model on the train set.
3. Generate `n\_pert\_per\_split` independent perturbed test sets with `perturb\_testset`, using the \*same\* DRO metric (`dro\_params`) but possibly different model (`method`).
4. On each perturbed test set, compute metrics via `metric\_fun`.
5. Aggregate metrics over perturbations into:
  - average performance (mean AUC / ACE),
  - worst-case performance (min AUC / max ACE).

Note (important):

- Even for `method = :logistic`, `dro\_params` is still required: it defines the ground metric that drives the test-set perturbations.
- DRO vs logistic are compared under the same \*shift model\*.

Returns:

- `Vector{NamedTuple}` , one per split, with fields:

```

        (split, auc_avg, auc_min, ace_avg, ace_max)
"""
function run_experiment(
    X::AbstractMatrix,
    Z::AbstractMatrix{<:Integer},
    group::AbstractVector{<:Integer},
    y::AbstractVector,
    ;
    train_ratio::Float64 = 0.7,
    method::Symbol = :dro,
    dro_params::Union{DROParams, Nothing} = nothing,
    logit_params::Union{LogitParams, Nothing} = nothing,
    pert_params::PerturbParams,
    n_splits::Int = 1,
    n_pert_per_split::Int = 100,
    metric_fun = default_metrics,
    rng = Random.default_rng(),
    optimizer = optimizer_with_attributes(Ipopt.Optimizer, "print_level"
)
    N = size(X, 1)
    @assert 0.0 < train_ratio < 1.0 "train_ratio must be in (0,1)"
    @assert size(Z, 1) == N
)

```

```

@assert length(group) == N
@assert length(y) == N
@assert dro_params !== nothing "dro_params is required (also for :logistic)"

dro = dro_params::DROParams

# Encode Z and group once, so DRO and logistic share the same encoding
Z_enc_all, G_enc_all, encinfo = encode_zg_reduced(Z, group)

results = NamedTuple[]

for s in 1:n_splits
    # 1. Random train/test split
    idx      = randperm(rng, N)
    n_train  = max(1, min(N - 1, round(Int, train_ratio * N)))
    train_idx = idx[1:n_train]
    test_idx  = idx[n_train+1:end]

    Xtr = X[train_idx, :]
    Ztr = Z[train_idx, :]
    gtr = group[train_idx]
    ytr = y[train_idx]

    Xte = X[test_idx, :]
    Zte = Z[test_idx, :]
    gte = group[test_idx]
    yte = y[test_idx]

    Zenc_tr = Z_enc_all[train_idx, :]
    Genc_tr = G_enc_all[train_idx, :]

    # 2. Fit model
    β =
        if method == :dro
            fit_dro_model(
                Xtr, Ztr, gtr, ytr,
                encinfo, dro;
                optimizer = optimizer,
            )
        elseif method == :logistic
            @assert logit_params !== nothing "logit_params must be provided"
            fit_logistic_model(
                Xtr, Zenc_tr, Genc_tr, ytr,
                encinfo, logit_params;
                optimizer = optimizer,
            )
        else
            error("Method $(method) not implemented. Use :dro or :logistic")
        end

    # 3. Monte Carlo over perturbed test sets
    metric_vals = Vector{NamedTuple}(undef, n_pert_per_split)

    for r in 1:n_pert_per_split
        Ū, Ź, Ÿ, Ÿ̂ = perturb_testset(
            Xte, Zte, gte, yte,
            encinfo, dro, pert_params;
            rng = rng,
        )
        probs = predict_proba(β, Ū, Ź, Ÿ)
    end
end

```

```

        metric_vals[r] = metric_fun(probs, ŷ)
    end

    # 4. Aggregate to average / worst-case
    aucs = [mv.auc for mv in metric_vals]
    aces = [mv.ace for mv in metric_vals]

    push!(results, (
        split = s,
        auc_avg = mean(aucs),
        auc_min = minimum(aucs),
        ace_avg = mean(aces),
        ace_max = maximum(aces),
    ))
end

return results
end

```

run\_experiment

## synthetic data experiment

In [91]:

```
#####
# Synthetic mixed-feature dataset generator
# (3 groups, 2D continuous, 1 categorical with 3 levels)
#####

"""
    generate_synthetic_mixed_data(N; rng = Random.default_rng())
Generate a synthetic dataset with:
- 3 groups g ∈ {1,2,3},
- 2 continuous features x ∈ ℝ²,
- 1 categorical feature z ∈ {1,2,3},
- labels y ∈ {-1,+1} drawn from a *known* logistic model f_{β*}.

The true logistic model uses the *same reduced encoding convention*
as our DRO / logistic models:
```

$$\begin{aligned}f_{\beta*}(x, z, g) &= \beta_0 \\&\quad + \beta_x^T x \\&\quad + \beta_z^T \varphi_z(z) \\&\quad + \beta_{grp}^T \varphi_g(g),\end{aligned}$$

where:

- categorical z has 3 levels, with reduced dummies for {1,2} and level 3 as baseline;
- group g has 3 levels, with reduced dummies for {1,2} and group 3 as baseline.

We fix (for illustration):

$$\begin{aligned}\beta_0 &= 0.0 \\ \beta_x &= [1.0, -0.5] \\ \beta_z &= [0.8, -0.8] \quad \# z = 1 \text{ positive shift, } z = 2 \text{ negative shift, } \\ \beta_{grp} &= [-0.5, 0.5] \quad \# \text{group 1 is slightly negative, group 2 slight}\end{aligned}$$

**Data-generating process**

-----  
For each sample  $i = 1, \dots, N$ :

1. Draw group  $g_i$ :  
 $P(g=1)=0.5, P(g=2)=0.3, P(g=3)=0.2.$
2. Conditional on  $g_i$ , draw categorical  $z_i$ :  
  - if  $g=1$ :  $P(z=1,2,3) = (0.5, 0.3, 0.2)$
  - if  $g=2$ :  $P(z=1,2,3) = (0.2, 0.5, 0.3)$
  - if  $g=3$ :  $P(z=1,2,3) = (0.3, 0.2, 0.5)$
3. Draw continuous features  $x_i \in \mathbb{R}^2$ :  
  - base group means:  
 $\mu_1 = (0, 0),$   
 $\mu_2 = (1, 1),$   
 $\mu_3 = (-1, 1);$
  - group-specific standard deviations:  
 $\sigma_1 = 0.5, \sigma_2 = 1.0, \sigma_3 = 0.7;$
  - category-specific shifts:  
 $\Delta_1 = (0.5, 0.0),$   
 $\Delta_2 = (0.0, -0.5),$   
 $\Delta_3 = (0.0, 0.0).$
  - set  $\mu_{\{g,z\}} = \mu_g + \Delta_z$  and draw  
 $x_i[j] \sim \text{Normal}(\mu_{\{g,z\}}[j], \sigma_g^2)$ , independently for  $j=1,2$ .
4. Compute the “true” logit:  
 $\eta_i = \beta_0 + \beta_x^\top x_i + \beta_z^\top \phi_z(z_i) + \beta_{grp}^\top \phi_g(g_i),$   
then draw  
 $y_i \sim \text{Bernoulli}(\sigma(\eta_i)),$   
and map  $\{0,1\}$  to  $\{-1,+1\}$ .

**Outputs**

-----  
-  $X :: \text{Matrix}\{\text{Float64}\}$  ( $N \times 2$ ) continuous features  
-  $Z :: \text{Matrix}\{\text{Int}\}$  ( $N \times 1$ ) categorical feature (values 1,2  
-  $\text{group} :: \text{Vector}\{\text{Int}\}$  (length  $N$ ) group indices (1,2,3)  
-  $y :: \text{Vector}\{\text{Int}\}$  (length  $N$ ) labels in  $\{-1,+1\}$   
-  $\beta_{\text{true}} :: \text{NamedTuple}$  ( $\beta_0, \beta_x, \beta_z, \beta_{grp}$ ), using \*reduced encoding  
 $\beta_{\text{true}}.\beta_z$  has length 2 (for  $z=1,2$ ;  $z=3$  baseline)  
 $\beta_{\text{true}}.\beta_{grp}$  has length 2 (for  $g=1,2$ ;  $g=3$  baseline)  
.....

```
function generate_synthetic_mixed_data(
    N::Int;
    rng = Random.default_rng(),
)
# -----
# 0. Dimensions
# -----
num_g = 3           # number of groups
n_x   = 2           # continuous features
m     = 1           # categorical components
k1    = 3           # categories for the single categorical feature

# -----
# 1. True logistic parameters in reduced encoding
# -----
β0_true      = 0.0
```

```

β_x_true = [1.0, -0.5] # length 2

# categorical z: 3 levels → 2 reduced dummies
# z=1 : φ_z(z) = (1,0) → coefficient β_z[1] = +0.8
# z=2 : φ_z(z) = (0,1) → coefficient β_z[2] = -0.8
# z=3 : baseline → (0,0)
β_z_true = [0.8, -0.8] # length 2

# group g: 3 levels → 2 reduced dummies
# g=1 : φ_g(g) = (1,0) → coefficient β_grp[1] = -0.5
# g=2 : φ_g(g) = (0,1) → coefficient β_grp[2] = +0.5
# g=3 : baseline → (0,0)
β_grp_true = [-0.5, 0.5] # length 2

# -----
# 2. Group and categorical distributions
# -----
# Group probabilities: (0.5, 0.3, 0.2)
group_probs = [0.5, 0.3, 0.2]
group_dist = Categorical(group_probs)

# Conditional categorical probabilities P(z | g)
# rows: g = 1,2,3; columns: z = 1,2,3
z_probs = [
    0.5 0.3 0.2; # g = 1
    0.2 0.5 0.3; # g = 2
    0.3 0.2 0.5 # g = 3
]
z_dists = [Categorical(z_probs[g, :]) for g in 1:num_g]

# -----
# 3. Continuous feature means and std per group/category
# -----
# base means μ_g
μ_g = [
    [0.0, 0.0], # g = 1
    [1.0, 1.0], # g = 2
    [-1.0, 1.0], # g = 3
]
μ_g

# group-specific standard deviations
σ_g = [0.5, 1.0, 0.7]

# category-specific shifts Δ_z
Δ_z = [
    [0.5, 0.0], # z = 1
    [0.0, -0.5], # z = 2
    [0.0, 0.0], # z = 3
]
Δ_z

# -----
# 4. Allocate output arrays
# -----
X = zeros(Float64, N, n_x) # N × 2
Z = Matrix{Int}(undef, N, m) # N × 1
group = Vector{Int}(undef, N)
y = Vector{Int}(undef, N)

# -----
# 5. Generate samples

```

```

# -----
for i in 1:N
    # (a) sample group
    g_i = rand(rng, group_dist)
    group[i] = g_i

    # (b) conditional categorical
    z_i = rand(rng, z_dists[g_i])   # ∈ {1,2,3}
    Z[i, 1] = z_i

    # (c) continuous x given (g,z)
    μ_base = μ_g[g_i]
    σ      = σ_g[g_i]
    μ_gz   = [μ_base[1] + Δ_z[z_i][1],
               μ_base[2] + Δ_z[z_i][2]]

    X[i, 1] = rand(rng, Normal(μ_gz[1], σ))
    X[i, 2] = rand(rng, Normal(μ_gz[2], σ))

    # (d) true logit η_i using the same reduced encoding as our model
    η = β₀_true + β_x_true[1]*X[i,1] + β_x_true[2]*X[i,2]

    # categorical contribution
    if z_i == 1
        η += β_z_true[1]
    elseif z_i == 2
        η += β_z_true[2]
    end

    # group contribution
    if g_i == 1
        η += β_grp_true[1]
    elseif g_i == 2
        η += β_grp_true[2]
    end

    # (e) sample label y_i ∈ {-1,+1}
    p = 1.0 / (1.0 + exp(-η))   # σ(η)
    y[i] = rand(rng) < p ? 1 : -1
end

β_true = (
    β₀      = β₀_true,
    β_x    = β_x_true,
    β_z    = β_z_true,
    β_grp = β_grp_true,
)

return X, Z, group, y, β_true
end

```

generate\_synthetic\_mixed\_data (generic function with 1 method)

In [92]:

```

#####
# Scenario U / V / R for synthetic experiment
#####

# Metric scenarios for (A_g, B_g, C_g), g = 1,2,3.
# U = uniform      : all groups same robustness

```

```

# V = "vulnerable" : group 1 easy to move, 3 hardest
# R = reversed      : group 1 hardest, 3 easiest

struct ScenarioSpec
    name::Symbol
    A_group::Vector{Float64}
    B_group::Vector{Float64}
    C_group::Vector{Float64}
end

const SCENARIO_U = ScenarioSpec(
    :U,
    [1.0, 1.0, 1.0],    # A_g
    [1.0, 1.0, 1.0],    # B_g
    [1.0, 1.0, 1.0],    # C_g
)

const SCENARIO_V = ScenarioSpec(
    :V,
    [0.5, 1.0, 2.0],    # group 1 cheap, group 3 expensive
    [0.5, 1.0, 2.0],
    [0.5, 1.0, 2.0],
)

const SCENARIO_R = ScenarioSpec(
    :R,
    [2.0, 1.0, 0.5],    # reverse of V
    [2.0, 1.0, 0.5],
    [2.0, 1.0, 0.5],
)

const SCENARIOS_UVR = [SCENARIO_U, SCENARIO_V, SCENARIO_R]

#####
# Build DROParams for the synthetic setting
#####

#####
# build_synth_dro_params(scen; theta=0.75)

Construct DROParams for the synthetic 2D+1cat, 3-group model
given a metric scenario `scen` and robustness level `theta`.

- gamma_x = [1, 1]
- delta    = [1]
- epsilon = -log(theta)
#####
function build_synth_dro_params(
    scen::ScenarioSpec;
    theta::Float64 = 0.75,
)
    n_x = 2    # 2 continuous features
    m   = 1    # 1 categorical component

    gamma_x = ones(Float64, n_x)
    delta   = ones(Float64, m)

```

```

epsilon = -log(theta)

return DROParams(
    delta,
    scen.A_group,
    scen.B_group,
    scen.C_group,
    gamma_x,
    epsilon,
)
end

#####
# UVR synthetic experiment driver (unpaired)
#####

using Random
using Statistics
using Ipopt
using JuMP

#####
run_synthetic_UVR();
    N = 3000,
    thetas = [0.5, 0.75, 0.9],
    w = 1.0,
    n_splits = 5,
    n_pert_per_split = 100,
    lambda_l2 = 0.01,
    seed = 2025,
    optimizer = optimizer_with_attributes(Ipopt.Optimizer, "print_lev
)

```

Run the synthetic experiment for scenarios U/V/R and both models (:dro, :logistic).

For each scenario  $s \in \{U, V, R\}$ , each  $\theta$ , and each method:

- Generate a fresh dataset of size  $N$ .
- Random 70/30 train/test split.
- Build DROParams using the scenario +  $\theta$ .
- Fit DRO or logistic model on training data.
- For each of `n\_pert\_per\_split` draws:
  - \* perturb the test set using the metric-based `perturb\_testset` with global noise scale  $w$ ;
  - \* compute metrics via `default\_metrics`.
- Aggregate metrics over perturbations and over `n\_splits`.

Returns a Vector of NamedTuples with fields:

```
(scenario, theta, method, auc_avg_mean, auc_min_mean,
 ace_avg_mean, ace_max_mean)
```

#####

```
function run_synthetic_UVR();
    N::Int = 3000,
    thetas::Vector{Float64} = [0.5, 0.75, 0.9],
    w::Float64 = 1.0,
    n_splits::Int = 5,
    n_pert_per_split::Int = 100,
    lambda_l2::Float64 = 0.01,
    seed::Int = 2025,
```

```

optimizer = optimizer_with_attributes(Iopt.Optimizer, "print_level"
)
rng = MersenneTwister(seed)

results = NamedTuple[]

for scen in SCENARIOS_UVR
    for theta in thetas
        # metric / DRO hyper-params (same for both models, used in pe
        dro_params = build_synth_dro_params(scen; theta = theta)
        pert_params = PerturbParams(:synthetic, w)
        logit_params = LogitParams(lambda_l2)

        for method in (:dro, :logistic)
            # store per-split summaries
            split_metrics = NamedTuple[]

            for s in 1:n_splits
                #
                # 1) Generate data + encode
                #
                X, Z, group, y, _ = generate_synthetic_mixed_data(N;

                Ntot = size(X, 1)
                idx = randperm(rng, Ntot)
                n_tr = max(1, min(Ntot - 1, round(Int, 0.7 * Ntot)))
                tr_idx = idx[1:n_tr]
                te_idx = idx[n_tr+1:end]

                Xtr = X[tr_idx, :]
                Ztr = Z[tr_idx, :]
                gtr = group[tr_idx]
                ytr = y[tr_idx]

                Xte = X[te_idx, :]
                Zte = Z[te_idx, :]
                gte = group[te_idx]
                yte = y[te_idx]

                Z_enc_all, G_enc_all, encinfo = encode_zg_reduced(Z,
                Zenc_tr = Z_enc_all[tr_idx, :]
                Genc_tr = G_enc_all[tr_idx, :]

                #
                # 2) Fit model
                #
                β =
                    if method == :dro
                        fit_dro_model(
                            Xtr, Ztr, gtr, ytr,
                            encinfo, dro_params;
                            optimizer = optimizer,
                        )
                    else
                        fit_logistic_model(
                            Xtr, Zenc_tr, Genc_tr, ytr,
                            encinfo, logit_params;
                            optimizer = optimizer,
                        )
                    end
            end
        end
    end
end

```

```

# -----
# 3) Monte Carlo over perturbed test sets
# -----
mvals = Vector{NamedTuple}(undef, n_pert_per_split)
for r in 1:n_pert_per_split
    X̃, Ź, ġ, ŷ = perturb_testset(
        Xte, Zte, gte, yte,
        encinfo, dro_params, pert_params;
        rng = rng,
    )
    probs = predict_proba(β, X̃, Ź, ġ)
    mvals[r] = default_metrics(probs, ŷ)
end

aucs = [mv.auc for mv in mvals]
aces = [mv.ace for mv in mvals]

push!(split_metrics, (
    auc_avg = mean(aucs),
    auc_min = minimum(aucs),
    ace_avg = mean(aces),
    ace_max = maximum(aces),
))
end

# average over splits
auc_avg_mean = mean(getfield.(split_metrics, :auc_avg))
auc_min_mean = mean(getfield.(split_metrics, :auc_min))
ace_avg_mean = mean(getfield.(split_metrics, :ace_avg))
ace_max_mean = mean(getfield.(split_metrics, :ace_max))

push!(results, (
    scenario = scen.name,
    theta = theta,
    method = method,
    auc_avg_mean = auc_avg_mean,
    auc_min_mean = auc_min_mean,
    ace_avg_mean = ace_avg_mean,
    ace_max_mean = ace_max_mean,
))
end
end
end
return results
end

```

run\_synthetic\_UVR

```
In [93]: res_uvr = run_synthetic_UVR(
    N = 3000,
    thetas = [0.5, 0.75, 0.9],
    w = 1.0,
    n_splits = 3,
    n_pert_per_split = 100,
    lambda_l2 = 0.01,
    seed = 2025,
)
foreach println, res_uvr)
```

```
(scenario = :U, theta = 0.5, method = :dro, auc_avg_mean = 0.6369017126449
766, auc_min_mean = 0.5982223959083725, ace_avg_mean = 0.1052586840305108
6, ace_max_mean = 0.13648512982352676)
(scenario = :U, theta = 0.5, method = :logistic, auc_avg_mean = 0.67782414
81233609, auc_min_mean = 0.646467419975348, ace_avg_mean = 0.1064596626861
5104, ace_max_mean = 0.13705157976967552)
(scenario = :U, theta = 0.75, method = :dro, auc_avg_mean = 0.643134442169
4218, auc_min_mean = 0.6003904143220631, ace_avg_mean = 0.0544526856574553
34, ace_max_mean = 0.08504948145570464)
(scenario = :U, theta = 0.75, method = :logistic, auc_avg_mean = 0.6795266
58752378, auc_min_mean = 0.6450369959346806, ace_avg_mean = 0.100066682452
58081, ace_max_mean = 0.12841126480213508)
(scenario = :U, theta = 0.9, method = :dro, auc_avg_mean = 0.6738718272725
711, auc_min_mean = 0.6376848680042937, ace_avg_mean = 0.1217474648982071
5, ace_max_mean = 0.14979496785620589)
(scenario = :U, theta = 0.9, method = :logistic, auc_avg_mean = 0.67990188
42579884, auc_min_mean = 0.6485459954142211, ace_avg_mean = 0.105806317472
30236, ace_max_mean = 0.13676428970281215)
(scenario = :V, theta = 0.5, method = :dro, auc_avg_mean = 0.6489033043393
408, auc_min_mean = 0.6126344510490052, ace_avg_mean = 0.1176209680086084
8, ace_max_mean = 0.14465760741437364)
(scenario = :V, theta = 0.5, method = :logistic, auc_avg_mean = 0.68176057
0456975, auc_min_mean = 0.6520507755459092, ace_avg_mean = 0.1089024945960
9053, ace_max_mean = 0.13765093246668003)
(scenario = :V, theta = 0.75, method = :dro, auc_avg_mean = 0.652476438175
8364, auc_min_mean = 0.6127705087033446, ace_avg_mean = 0.1637115819556284
3, ace_max_mean = 0.19408163063056727)
(scenario = :V, theta = 0.75, method = :logistic, auc_avg_mean = 0.6690869
944919505, auc_min_mean = 0.6338598393608436, ace_avg_mean = 0.11547342307
60941, ace_max_mean = 0.1449506996548765)
(scenario = :V, theta = 0.9, method = :dro, auc_avg_mean = 0.6508075290759
426, auc_min_mean = 0.6195720811114772, ace_avg_mean = 0.2612878281711342,
ace_max_mean = 0.2954031485307242)
(scenario = :V, theta = 0.9, method = :logistic, auc_avg_mean = 0.66373709
52903664, auc_min_mean = 0.6305225439162515, ace_avg_mean = 0.124598500465
38871, ace_max_mean = 0.15221074707276644)
(scenario = :R, theta = 0.5, method = :dro, auc_avg_mean = 0.6478109398462
201, auc_min_mean = 0.6169224657208989, ace_avg_mean = 0.1124988843627217
1, ace_max_mean = 0.1406319234818605)
(scenario = :R, theta = 0.5, method = :logistic, auc_avg_mean = 0.69531659
6369218, auc_min_mean = 0.6552489433493119, ace_avg_mean = 0.0869355273722
3868, ace_max_mean = 0.12527254933759782)
(scenario = :R, theta = 0.75, method = :dro, auc_avg_mean = 0.631805362391
2702, auc_min_mean = 0.593463122354056, ace_avg_mean = 0.1015542186644509
1, ace_max_mean = 0.13764007608361709)
(scenario = :R, theta = 0.75, method = :logistic, auc_avg_mean = 0.7077750
451702617, auc_min_mean = 0.6773097190912557, ace_avg_mean = 0.07687866114
552354, ace_max_mean = 0.10425795074675649)
(scenario = :R, theta = 0.9, method = :dro, auc_avg_mean = 0.6585341836979
524, auc_min_mean = 0.623625140882094, ace_avg_mean = 0.13056722453192135,
ace_max_mean = 0.1559706987798008)
(scenario = :R, theta = 0.9, method = :logistic, auc_avg_mean = 0.70062344
51115724, auc_min_mean = 0.6644491733075465, ace_avg_mean = 0.077460024677
60983, ace_max_mean = 0.10477428508923033)
```

## Real-World Data

In [69]: **using** Statistics

```

summarize_experiment(res_vec)

Given the vector of per-split results returned by `run_experiment`,
compute simple means across splits.

Each element of `res_vec` is expected to have fields:
    (split, auc_avg, auc_min, ace_avg, ace_max)

Returns a NamedTuple:
    (auc_avg_mean, auc_min_mean, ace_avg_mean, ace_max_mean)
"""

function summarize_experiment(res_vec::Vector{<:NamedTuple})
    auc_avg_mean = mean(r.auc_avg for r in res_vec)
    auc_min_mean = mean(r.auc_min for r in res_vec)
    ace_avg_mean = mean(r.ace_avg for r in res_vec)
    ace_max_mean = mean(r.ace_max for r in res_vec)
    return (
        auc_avg_mean = auc_avg_mean,
        auc_min_mean = auc_min_mean,
        ace_avg_mean = ace_avg_mean,
        ace_max_mean = ace_max_mean,
    )
end

```

summarize\_experiment

```

In [73]: using Random
using Ipopt
using CSV, DataFrames # harmless if already loaded

"""

run_churn_scenarios(
    path;
    severities = [:mild, :strong],
    thetas = [0.5, 0.75, 0.9],
    train_ratio = 0.7,
    n_splits = 2,
    n_pert_per_split = 100,
    lambda_l2 = 0.01,
    seed = 2025,
    max_samples_per_scenario = nothing,
    optimizer = optimizer_with_attributes(Ipopt.Optimizer, "print_lev
)

```

Run subgroup-DRO vs standard logistic regression on the Bank Customer Churn dataset, under a grid of (severity,  $\theta$ ) settings.

For each (severity,  $\theta$ ):

1. Optionally subsample at most `max\_samples\_per\_scenario` points.
2. Build DROParams and PerturbParams via `build\_churn\_params`.
3. Run `run\_experiment` with method = :dro.
4. Run `run\_experiment` with method = :logistic.
5. Summarize each via `summarize\_experiment`.

Inputs

-----

- path	: path to "Bank Customer Churn Prediction.csv".
- severities	: e.g. [:mild, :strong].
- thetas	: robustness levels $\theta$ in (0,1).
- train_ratio	: fraction of samples used for training in each

```

- n_splits : number of random train/test splits.
- n_pert_per_split : number of perturbed test sets per split.
- lambda_l2 : L2-regularization coefficient for logistic ba
- seed : base random seed.
- max_samples_per_scenario: if not `nothing`, limit each (severity, θ) sc
                           to this many samples by random subsampling.
- optimizer : JuMP optimizer for both models.

>Returns
-----
- results :: Vector{NamedTuple}

Each element has fields:
  (severity, theta, method,
   auc_avg_mean, auc_min_mean, ace_avg_mean, ace_max_mean)
.....
function run_churn_scenarios(
    path::AbstractString;
    severities = [:mild, :strong],
    thetas = [0.5, 0.75, 0.9],
    train_ratio::Float64 = 0.7,
    n_splits::Int = 2,
    n_pert_per_split::Int = 100,
    lambda_l2::Float64 = 0.01,
    seed::Int = 2025,
    max_samples_per_scenario::Union{Nothing, Int} = nothing,
    optimizer = optimizer_with_attributes(Ipopt.Optimizer, "print_level"
)
    # Load the full dataset once
    X_full, Z_full, g_full, y_full, meta = load_churn_dataset(path)
    N_full = size(X_full, 1)

    # Logistic baseline hyperparameter
    logit_params = LogitParams(lambda_l2)

    # One RNG is enough here; runs are unpaired but share the same distri
    rng = Random.MersenneTwister(seed)

    results = NamedTuple[]

    # Loop over (severity, theta)
    for sev in severities
        for θ in thetas
            # Decide which subset of samples to use for this (sev, θ) sce
            if max_samples_per_scenario === nothing || max_samples_per_sc
                # Use all samples
                X = X_full
                Z = Z_full
                g = g_full
                y = y_full
            else
                # Randomly choose a subset of size max_samples_per_scenar
                idx_sub = randperm(rng, N_full)[1:max_samples_per_scenari
                X = X_full[idx_sub, :]
                Z = Z_full[idx_sub, :]
                g = g_full[idx_sub]
                y = y_full[idx_sub]
            end
            n_x  = size(X, 2)
        end
    end
end

```

```

m      = size(Z, 2)
num_g = maximum(g)

# Build DRO and perturbation parameters for this scenario
dro_params, pert_params =
    build_churn_params(sev, θ; n_x = n_x, m = m, num_g = num_g)

# Run DRO model on this (possibly subsampled) dataset
res_dro = run_experiment(
    X, Z, g, y;
    train_ratio      = train_ratio,
    method           = :dro,
    dro_params       = dro_params,
    pert_params      = pert_params,
    n_splits         = n_splits,
    n_pert_per_split = n_pert_per_split,
    rng              = rng,
    optimizer        = optimizer,
)
sum_dro = summarize_experiment(res_dro)

push!(results, (
    severity      = sev,
    theta         = θ,
    method        = :dro,
    auc_avg_mean = sum_dro.auc_avg_mean,
    auc_min_mean = sum_dro.auc_min_mean,
    ace_avg_mean = sum_dro.ace_avg_mean,
    ace_max_mean = sum_dro.ace_max_mean,
))

# Run logistic baseline on the same (possibly subsampled) dat
res_log = run_experiment(
    X, Z, g, y;
    train_ratio      = train_ratio,
    method           = :logistic,
    logit_params     = logit_params,
    pert_params      = pert_params,
    n_splits         = n_splits,
    n_pert_per_split = n_pert_per_split,
    rng              = rng,
    optimizer        = optimizer,
)
sum_log = summarize_experiment(res_log)

push!(results, (
    severity      = sev,
    theta         = θ,
    method        = :logistic,
    auc_avg_mean = sum_log.auc_avg_mean,
    auc_min_mean = sum_log.auc_min_mean,
    ace_avg_mean = sum_log.ace_avg_mean,
    ace_max_mean = sum_log.ace_max_mean,
))

end
end

return results
end

```

## run\_churn\_scenarios

```
In [75]: path = "Bank Customer Churn Prediction.csv"
```

```
results_churn = run_churn_scenarios(  
    path;  
    severities = [:mild, :strong],  
    thetas = [0.5, 0.75, 0.9],  
    train_ratio = 0.7,  
    n_splits = 8,  
    n_pert_per_split = 100,  
    lambda_l2 = 0.01,  
    seed = 2025,  
    max_samples_per_scenario = 2000,    # <-- this is the cap  
)  
  
println.(results_churn)
```

```
(severity = :mild, theta = 0.5, method = :dro, auc_avg_mean = 0.6809890560  
473452, auc_min_mean = 0.664903204398336, ace_avg_mean = 0.064601779854896  
18, ace_max_mean = 0.08777094974627037)  
(severity = :mild, theta = 0.5, method = :logistic, auc_avg_mean = 0.75178  
45869875566, auc_min_mean = 0.7278152044312919, ace_avg_mean = 0.041937274  
50301925, ace_max_mean = 0.0660311427357998)  
(severity = :mild, theta = 0.75, method = :dro, auc_avg_mean = 0.730040657  
7070826, auc_min_mean = 0.7116409109201876, ace_avg_mean = 0.0537579125904  
25276, ace_max_mean = 0.07300932894970448)  
(severity = :mild, theta = 0.75, method = :logistic, auc_avg_mean = 0.7333  
472257349112, auc_min_mean = 0.7081167983791491, ace_avg_mean = 0.03871367  
7706732764, ace_max_mean = 0.06103026321394643)  
(severity = :mild, theta = 0.9, method = :dro, auc_avg_mean = 0.7527691053  
883292, auc_min_mean = 0.7287964524184268, ace_avg_mean = 0.06068306953020  
5413, ace_max_mean = 0.08226691094058226)  
(severity = :mild, theta = 0.9, method = :logistic, auc_avg_mean = 0.74358  
17374996072, auc_min_mean = 0.7174550020938562, ace_avg_mean = 0.042416671  
69189757, ace_max_mean = 0.06407911341649271)  
(severity = :strong, theta = 0.5, method = :dro, auc_avg_mean = 0.64121508  
8869021, auc_min_mean = 0.616288977723691, ace_avg_mean = 0.07629882130515  
442, ace_max_mean = 0.09647496631886289)  
(severity = :strong, theta = 0.5, method = :logistic, auc_avg_mean = 0.709  
687278929923, auc_min_mean = 0.6693135594425055, ace_avg_mean = 0.04870220  
426575063, ace_max_mean = 0.07183562847552016)  
(severity = :strong, theta = 0.75, method = :dro, auc_avg_mean = 0.6876614  
974775004, auc_min_mean = 0.6615166921501202, ace_avg_mean = 0.05839039397  
970805, ace_max_mean = 0.07868046922392939)  
(severity = :strong, theta = 0.75, method = :logistic, auc_avg_mean = 0.69  
26513102271572, auc_min_mean = 0.6490075944510951, ace_avg_mean = 0.048570  
312242927335, ace_max_mean = 0.07465620969977059)  
(severity = :strong, theta = 0.9, method = :dro, auc_avg_mean = 0.70652008  
17874847, auc_min_mean = 0.6754661556606002, ace_avg_mean = 0.044212101836  
086294, ace_max_mean = 0.06705531415980445)  
(severity = :strong, theta = 0.9, method = :logistic, auc_avg_mean = 0.705  
4237888163174, auc_min_mean = 0.6677568197490297, ace_avg_mean = 0.0442506  
4308727014, ace_max_mean = 0.06873497159443204)
```

```
12-element Vector{Nothing}:
nothing
```

```
In [76]: using DataFrames

"""

    summarize_differences(results_churn)

Convert the vector of NamedTuples returned by `run_churn_scenarios` into a DataFrame with one row per (severity, theta), and columns:

- auc_dro, auc_log, delta_auc = auc_dro - auc_log
- ace_dro, ace_log, delta_ace = ace_dro - ace_log
"""

function summarize_differences(results_churn)
    df = DataFrame(results_churn)

    # pivot: two rows per config -> one row per (severity, theta)
    configs = unique(df[:, [:severity, :theta]])
    out = DataFrame(
        severity = String[],
        theta = Float64[],
        auc_dro = Float64[],
        auc_log = Float64[],
        delta_auc = Float64[],
        ace_dro = Float64[],
        ace_log = Float64[],
        delta_ace = Float64[],
    )

    for row in eachrow(configs)
        sev = row.severity
        θ = row.theta

        sub = df[(df.severity .== sev) .& (df.theta .== θ), :]

        dro_row = sub[sub.method .== :dro, :]
        log_row = sub[sub.method .== :logistic, :]

        auc_dro = dro_row.auc_avg_mean[1]
        auc_log = log_row.auc_avg_mean[1]
        ace_dro = dro_row.ace_avg_mean[1]
        ace_log = log_row.ace_avg_mean[1]

        push!(out, (
            string(sev),
            θ,
            auc_dro,
            auc_log,
```

```

        auc_dro - auc_log,
        ace_dro,
        ace_log,
        ace_dro - ace_log,
    ))
end

return out
end

df_diff = summarize_differences(results_churn)
println(df_diff)

```

Row	severity	theta	auc_dro	auc_log	delta_auc	ace_dro	ace_
log		delta_ace					
t64	String	Float64	Float64	Float64	Float64	Float64	Floa
		Float64					
1	mild	0.5	0.680989	0.751785	-0.0707955	0.0646018	0.04
19373	<i>0.0226645</i>						
2	mild	0.75	0.730041	0.733347	-0.00330657	0.0537579	0.03
87137	<i>0.0150442</i>						
3	mild	0.9	0.752769	0.743582	0.00918737	0.0606831	0.04
24167	<i>0.0182664</i>						
4	strong	0.5	0.641215	0.709687	-0.0684722	0.0762988	0.04
87022	<i>0.0275966</i>						
5	strong	0.75	0.687661	0.692651	-0.00498981	0.0583904	0.04
85703	<i>0.00982008</i>						
6	strong	0.9	0.70652	0.705424	0.00109629	0.0442121	0.04
42506	<i>-3.85413e-5</i>						

In [95]:

```

#####
# Churn: scenario-based DRO + perturb parameters (U / V / R)
#####

"""
    build_churn_scenario_params(scenario, theta;
                                n_x, m, num_g, w = 1.0)

Build (dro_params, pert_params) for the Bank Churn dataset
under one of three subgroup metric scenarios:

- :U (Uniform): A_g = B_g = C_g = 1           for all groups.
- :V (Vulnerable): (A,B,C) = (0.5,0.5,0.5) for group 1,
                     (1,1,1)                  for group 2,
                     (2,2,2)                  for group 3.
- :R (Reversed): same pattern as :V but swap group 1 and 3.

Here we keep the data-generating side (γ_x, δ) simple:

- γ_x[j] = 1      for all continuous features (already standardized).
- δ[ℓ] = 1       for all categorical features.
- ε = -log(theta).

The perturbation object only carries (mode, w); the *metric*
(A_g · γ_x, B_g · δ, C_g) controls how large the actual shifts are
inside `perturb_testset`.
"""

```

```

function build_churn_scenario_params(
    scenario::Symbol,
    theta::Float64;
    n_x::Int,
    m::Int,
    num_g::Int,
    w::Float64 = 1.0,
)
    @assert num_g == 3 "This helper assumes 3 groups (France / Germany /"

    # Base feature scales
    gamma_x = ones(Float64, n_x)
    delta = ones(Float64, m)

    # Scenario-specific A_g, B_g, C_g
    A_group = zeros(Float64, num_g)
    B_group = zeros(Float64, num_g)
    C_group = zeros(Float64, num_g)

    if scenario == :U
        # All groups equally easy to move
        A_group .= 1.0
        B_group .= 1.0
        C_group .= 1.0

    elseif scenario == :V
        # Group 1 most vulnerable, group 3 most robust
        A_group .= [0.5, 1.0, 2.0]
        B_group .= [0.5, 1.0, 2.0]
        C_group .= [0.5, 1.0, 2.0]

    elseif scenario == :R
        # Reverse: group 3 most vulnerable, group 1 most robust
        A_group .= [2.0, 1.0, 0.5]
        B_group .= [2.0, 1.0, 0.5]
        C_group .= [2.0, 1.0, 0.5]

    else
        error("Unknown scenario = $scenario. Use :U, :V or :R.")
    end

    # DRO radius
    epsilon = -log(theta)

    dro_params = DROParams(
        delta,
        A_group,
        B_group,
        C_group,
        gamma_x,
        epsilon,
    )

    # Perturbation knob w: global noise budget
    pert_params = PerturbParams(scenario, w)

    return dro_params, pert_params
end

```

build\_churn\_scenario\_params

In [100...]

```
#####
# Churn experiment driver (U / V / R, 1000-sample subset)
#####

using Random
using Statistics
using Ipopt

#####
run_churn_scenarios_uvr(
    path;
    scenarios = [:U, :V, :R],
    thetas = [0.5, 0.75, 0.9],
    train_ratio = 0.7,
    n_splits = 5,
    n_pert_per_split = 100,
    lambda_l2 = 1e-2,
    max_samples = 1000,
    w = 1.0,
    seed = 2025,
    optimizer = optimizer_with_attributes(Ipopt.Optimizer, "print_level"
)
```

Apply the same U / V / R subgroup-metric pipeline to the Bank Customer Churn dataset, using only a random subset of at most `max\_samples` points for speed.

For each scenario  $\in \{U, V, R\}$  and  $\theta \in \{0.5, 0.75, 0.9\}$ :

1. Build (dro\_params, pert\_params) via `build\_churn\_scenario\_params`.
2. Run `run\_experiment` with method = :dro.
3. Run `run\_experiment` with method = :logistic (same perturbation law).
4. Aggregate per-split summaries into means over splits.

Returns a vector of NamedTuples with fields:

```
(scenario, theta, method, auc_avg_mean, auc_min_mean,
ace_avg_mean, ace_max_mean)
```

```
#####
function run_churn_scenarios_uvr(
    path::AbstractString;
    scenarios = [:U, :V, :R],
    thetas = [0.5, 0.75, 0.9],
    train_ratio::Float64 = 0.7,
    n_splits::Int = 5,
    n_pert_per_split::Int = 100,
    lambda_l2::Float64 = 1e-2,
    max_samples::Int = 1000,
    w::Float64 = 1.0,
    seed::Int = 2025,
    optimizer = optimizer_with_attributes(Ipopt.Optimizer, "print_level"
)
    rng = MersenneTwister(seed)

    # 1) Load and subsample dataset
    X_full, Z_full, g_full, y_full, meta = load_churn_dataset(path; stand
    N_total = size(X_full, 1)
    N_use = min(max_samples, N_total)

    idx_sub = randperm(rng, N_total)[1:N_use]
```

```

X = X_full[idx_sub, :]
Z = Z_full[idx_sub, :]
g = g_full[idx_sub]
y = y_full[idx_sub]

n_x = size(X, 2)
m = size(Z, 2)
num_g = maximum(g)

logit_params = LogitParams(lambda_l2)

results = NamedTuple[]

for scen in scenarios
    println(scen)
    for θ in thetas
        println(θ)
        # 2) Build DRO + perturbation parameters for this scenario
        dro_params, pert_params = build_churn_scenario_params(
            scen, θ;
            n_x = n_x,
            m = m,
            num_g = num_g,
            w = w,
        )

        # 3) DRO model
        res_dro = run_experiment(
            X, Z, g, y;
            train_ratio = train_ratio,
            method = :dro,
            dro_params = dro_params,
            logit_params = logit_params,
            pert_params = pert_params,
            n_splits = n_splits,
            n_pert_per_split = n_pert_per_split,
            rng = rng,
            optimizer = optimizer,
        )

        # 4) Logistic baseline
        res_logit = run_experiment(
            X, Z, g, y;
            train_ratio = train_ratio,
            method = :logistic,
            dro_params = dro_params,
            logit_params = logit_params,
            pert_params = pert_params,
            n_splits = n_splits,
            n_pert_per_split = n_pert_per_split,
            rng = rng,
            optimizer = optimizer,
        )

        # 5) Average the per-split summaries
        for (method, res) in zip(:dro, :logistic), (res_dro, res_logit)
            auc_avg_mean = mean(r.auc_avg for r in res)
            auc_min_mean = mean(r.auc_min for r in res)
            ace_avg_mean = mean(r.ace_avg for r in res)
            ace_max_mean = mean(r.ace_max for r in res)
        end
    end
end

```

```

        push!(results, (
            scenario      = scen,
            theta         = θ,
            method        = method,
            auc_avg_mean = auc_avg_mean,
            auc_min_mean = auc_min_mean,
            ace_avg_mean = ace_avg_mean,
            ace_max_mean = ace_max_mean,
        ))
    end
end

return results
end

```

run\_churn\_scenarios\_uvr

```
In [101]: path = "Bank Customer Churn Prediction.csv"

results_churn_uvr = run_churn_scenarios_uvr(
    path;
    scenarios = [:U, :V, :R],
    thetas = [0.5, 0.75, 0.9],
    train_ratio = 0.7,
    n_splits = 5,
    n_pert_per_split = 100,
    lambda_l2 = 1e-2,
    max_samples = 1000,
    w = 1.0,
    seed = 2025,
)
```

U  
0.5  
0.75  
0.9  
V  
0.5  
0.75  
0.9  
R  
0.5  
0.75  
0.9

```
18-element Vector{NamedTuple}:
  (scenario = :U, theta = 0.5, method = :dro, auc_avg_mean = 0.600672888496
  1647, auc_min_mean = 0.5372019670609568, ace_avg_mean = 0.0850202652264732
  1, ace_max_mean = 0.11759724874090113)
  (scenario = :U, theta = 0.5, method = :logistic, auc_avg_mean = 0.6218250
  865737012, auc_min_mean = 0.5501126897822066, ace_avg_mean = 0.09438246192
  760315, ace_max_mean = 0.131438654906637)
  (scenario = :U, theta = 0.75, method = :dro, auc_avg_mean = 0.58192791423
  89265, auc_min_mean = 0.5203608841970363, ace_avg_mean = 0.061506292966849
  45, ace_max_mean = 0.0980818336640321)
  (scenario = :U, theta = 0.75, method = :logistic, auc_avg_mean = 0.640034
  4725704978, auc_min_mean = 0.5679586614900688, ace_avg_mean = 0.0856502788
  3003404, ace_max_mean = 0.12288719969875483)
  (scenario = :U, theta = 0.9, method = :dro, auc_avg_mean = 0.608438816389
  667, auc_min_mean = 0.5313931089558744, ace_avg_mean = 0.0634777225748133
  6, ace_max_mean = 0.1044129528383747)
  (scenario = :U, theta = 0.9, method = :logistic, auc_avg_mean = 0.6289112
  7677067, auc_min_mean = 0.567954691822777, ace_avg_mean = 0.08911254339202
  858, ace_max_mean = 0.1321977006832086)
  (scenario = :V, theta = 0.5, method = :dro, auc_avg_mean = 0.588502759409
  3508, auc_min_mean = 0.5187536229290397, ace_avg_mean = 0.0733631312098519
  3, ace_max_mean = 0.10997329719691713)
  (scenario = :V, theta = 0.5, method = :logistic, auc_avg_mean = 0.6245864
  297487123, auc_min_mean = 0.5573429388070347, ace_avg_mean = 0.09567139063
  447101, ace_max_mean = 0.13412318329783363)
  (scenario = :V, theta = 0.75, method = :dro, auc_avg_mean = 0.59240070116
  84548, auc_min_mean = 0.5263368631427828, ace_avg_mean = 0.079699904265769
  71, ace_max_mean = 0.1167856060440502)
  (scenario = :V, theta = 0.75, method = :logistic, auc_avg_mean = 0.628711
  0445119499, auc_min_mean = 0.5583531649786219, ace_avg_mean = 0.0950099980
  3440499, ace_max_mean = 0.13367615562858454)
  (scenario = :V, theta = 0.9, method = :dro, auc_avg_mean = 0.572953241281
  928, auc_min_mean = 0.5074456172918075, ace_avg_mean = 0.0649898930305766
  6, ace_max_mean = 0.10801462505590768)
  (scenario = :V, theta = 0.9, method = :logistic, auc_avg_mean = 0.6141023
  333814848, auc_min_mean = 0.5388785482171677, ace_avg_mean = 0.09847893095
  235506, ace_max_mean = 0.14231754452958617)
  (scenario = :R, theta = 0.5, method = :dro, auc_avg_mean = 0.589469064482
  9353, auc_min_mean = 0.5249913056751174, ace_avg_mean = 0.0741379083302800
  4, ace_max_mean = 0.11579876071421318)
  (scenario = :R, theta = 0.5, method = :logistic, auc_avg_mean = 0.6323692
  667771349, auc_min_mean = 0.5556010820839283, ace_avg_mean = 0.08957496391
  661526, ace_max_mean = 0.13482479317835475)
  (scenario = :R, theta = 0.75, method = :dro, auc_avg_mean = 0.60625271859
  65222, auc_min_mean = 0.5313759589766944, ace_avg_mean = 0.077975361046867
  64, ace_max_mean = 0.10949653105602628)
  (scenario = :R, theta = 0.75, method = :logistic, auc_avg_mean = 0.633821
  4173352161, auc_min_mean = 0.5647782646969348, ace_avg_mean = 0.0913155834
  2989252, ace_max_mean = 0.13443504653581545)
  (scenario = :R, theta = 0.9, method = :dro, auc_avg_mean = 0.596405906774
  8367, auc_min_mean = 0.5253347857068538, ace_avg_mean = 0.0588723106391249
  24, ace_max_mean = 0.09415360015661693)
  (scenario = :R, theta = 0.9, method = :logistic, auc_avg_mean = 0.6264880
  687261054, auc_min_mean = 0.556192388761948, ace_avg_mean = 0.092351231718
  19284, ace_max_mean = 0.13546810519900426)
```

In [ ]: